



Planning and Scheduling: The Agents Metaphor



Hochschule
Bonn-Rhein-Sieg



Prof. Dr.-Ing. Gerhard K. Kraetzschmar


Acknowledgements

- These slides refer to Chapter 2 of the textbook:
S. Russell and P. Norvig:
Artificial Intelligence: A Modern Approach
Prentice Hall, 2003, 2nd Edition (or more recent edition)
- These slides are an adaptation of slides by Min-Yen Kan
- The contributions of these authors are gratefully acknowledged.

- Agents and environments
- Rationality
- PEAS
 - Performance measure
 - Environment
 - Actuators
 - Sensors
- Environment types
- Agent types

- An **agent** is anything that can be viewed as **perceiving** its **environment** through **sensors** and **acting** upon that environment through **actuators**
- Human agent:
 - eyes, ears, and other organs for sensors;
 - hands, legs, mouth, and other body parts for actuators
- Robotic agent:
 - cameras and infrared range finders for sensors;
 - various motors for actuators

Agents and environments




- The **agent function** maps from percept histories to actions:

$$f : P^* \mapsto A$$

- The **agent program** runs on the physical **architecture** to produce f
- $\text{agent} = \text{architecture} + \text{program}$

Vacuum-Cleaner World



- Percepts: location and contents, e.g., [A,Dirty]
- Actions: Left, Right, Suck, NoOp

A Vacuum-Cleaner Agent

Percept sequence	Action
$[A, \textit{Clean}]$	<i>Right</i>
$[A, \textit{Dirty}]$	<i>Suck</i>
$[B, \textit{Clean}]$	<i>Left</i>
$[B, \textit{Dirty}]$	<i>Suck</i>
$[A, \textit{Clean}], [A, \textit{Clean}]$	<i>Right</i>
$[A, \textit{Clean}], [A, \textit{Dirty}]$	<i>Suck</i>
\vdots	\vdots

function REFLEX-VACUUM-AGENT($[location, status]$) **returns** an action

if $status = \textit{Dirty}$ **then return** *Suck*
else if $location = A$ **then return** *Right*
else if $location = B$ **then return** *Left*

- An agent should strive to "do the right thing", based on what it can perceive and the actions it can perform. The right action will cause the agent to be most successful.
- Performance measure:
An objective criterion for success of an agent's behavior.
- Example: the performance measure of a vacuum-cleaner agent could be
 - the amount of dirt cleaned up,
 - the amount of time taken,
 - the amount of electricity consumed,
 - the amount of noise generated, etc.

Rational Agents

- **Rational** Agent: For each possible percept sequence, a rational agent should select an action that is expected to maximize its performance measure, given the evidence provided by the percept sequence and whatever built-in knowledge the agent has
- Rationality is distinct from omniscience (all-knowing with infinite knowledge)
- Agents can perform actions in order to modify future percepts so as to obtain useful information (information gathering, exploration)
- An agent is autonomous if its behavior is determined by its own experience (with ability to learn and adapt)

- We must first specify the setting for intelligent agent design
- This is done in terms of PEAS descriptions:
 - Performance measure
 - Environment
 - Actuators
 - Sensors
- Consider, e.g., the task of designing an automated taxi driver:
- Agent: taxi driver
 - Performance measure: safe, fast, legal, comfortable trip, maximize profits
 - Environment: roads, other traffic, pedestrians, customers
 - Actuators: steering wheel, accelerator, brake, signal, horn
 - Sensors: cameras, sonar, speedometer, GPS, odometer, engine sensors, keyboard

- Agent: medical diagnosis system
 - Performance measure: healthy patient, minimize costs, lawsuits
 - Environment: patient, hospital, staff
 - Actuators: screen display (questions, tests, diagnoses, treatments, referrals)
 - Sensors: keyboard (entry of symptoms, findings, patient's answers)
- Agent: part-picking robot
 - Performance measure: percentage of parts in correct bins
 - Environment: conveyor belt with parts, bins
 - Actuators: jointed arm and hand
 - Sensors: camera, joint angle sensors
- Agent: interactive English tutor
 - Performance measure: maximize student's score on test
 - Environment: set of students
 - Actuators: screen display (exercises, suggestions, corrections)
 - Sensors: keyboard

Environment Types

- **Fully observable** vs. **partially observable**:
 - Agent has access to the complete state of the environment at each point in time.
- **Deterministic** vs. **stochastic**:
 - The next state of the environment is completely determined by the current state and the action executed by the agent. (If the environment is deterministic except for the actions of other agents, then the environment is strategic.)
- **Episodic** vs. **sequential**:
 - The agent's experience is divided into atomic "episodes" (each episode consists of the agent perceiving and then performing a single action), and the choice of action in each episode depends only on the episode itself.

Environment Types

- **Static** (vs. dynamic):
 - The environment is unchanged while an agent is deliberating.
(The environment is semidynamic if the environment itself does not change with the passage of time but the agent's performance score does)
- **Discrete** (vs. continuous):
 - A limited number of distinct, clearly defined percepts and actions.
- **Single agent** (vs. multiagent):
 - An agent operating by itself in an environment.

Environment Types

	Chess w. clock	Chess w.o. clock	Taxi driving
■ Fully observable	Yes	Yes	No
■ Deterministic	Strategic	Strategic	No
■ Episodic	No	No	No
■ Static	Semi	Yes	No
■ Discrete	Yes	Yes	No
■ Single agent	No	No	No

- The environment type largely determines the agent design
- The real world is (of course) partially observable, stochastic, sequential, dynamic, continuous, and multiagent

Agent Functions and Programs

- An agent is completely specified by the **agent function** mapping percept sequences to actions
- One agent function (or a small equivalence class) is **rational**
- Aim: find a way to implement the rational agent function concisely

Table-Lookup Agent

Percept sequence	Action
<i>[A, Clean]</i>	<i>Right</i>
<i>[A, Dirty]</i>	<i>Suck</i>
<i>[B, Clean]</i>	<i>Left</i>
<i>[B, Dirty]</i>	<i>Suck</i>
<i>[A, Clean], [A, Clean]</i>	<i>Right</i>
<i>[A, Clean], [A, Dirty]</i>	<i>Suck</i>
<i>⋮</i>	<i>⋮</i>

- Drawbacks:
 - Huge table
 - Takes a long time to build the table
 - No autonomy
 - Even with learning, need a long time to learn the table entries

Agent Program For a Vacuum-Cleaner Agent

function REFLEX-VACUUM-AGENT([*location,status*]) returns an action

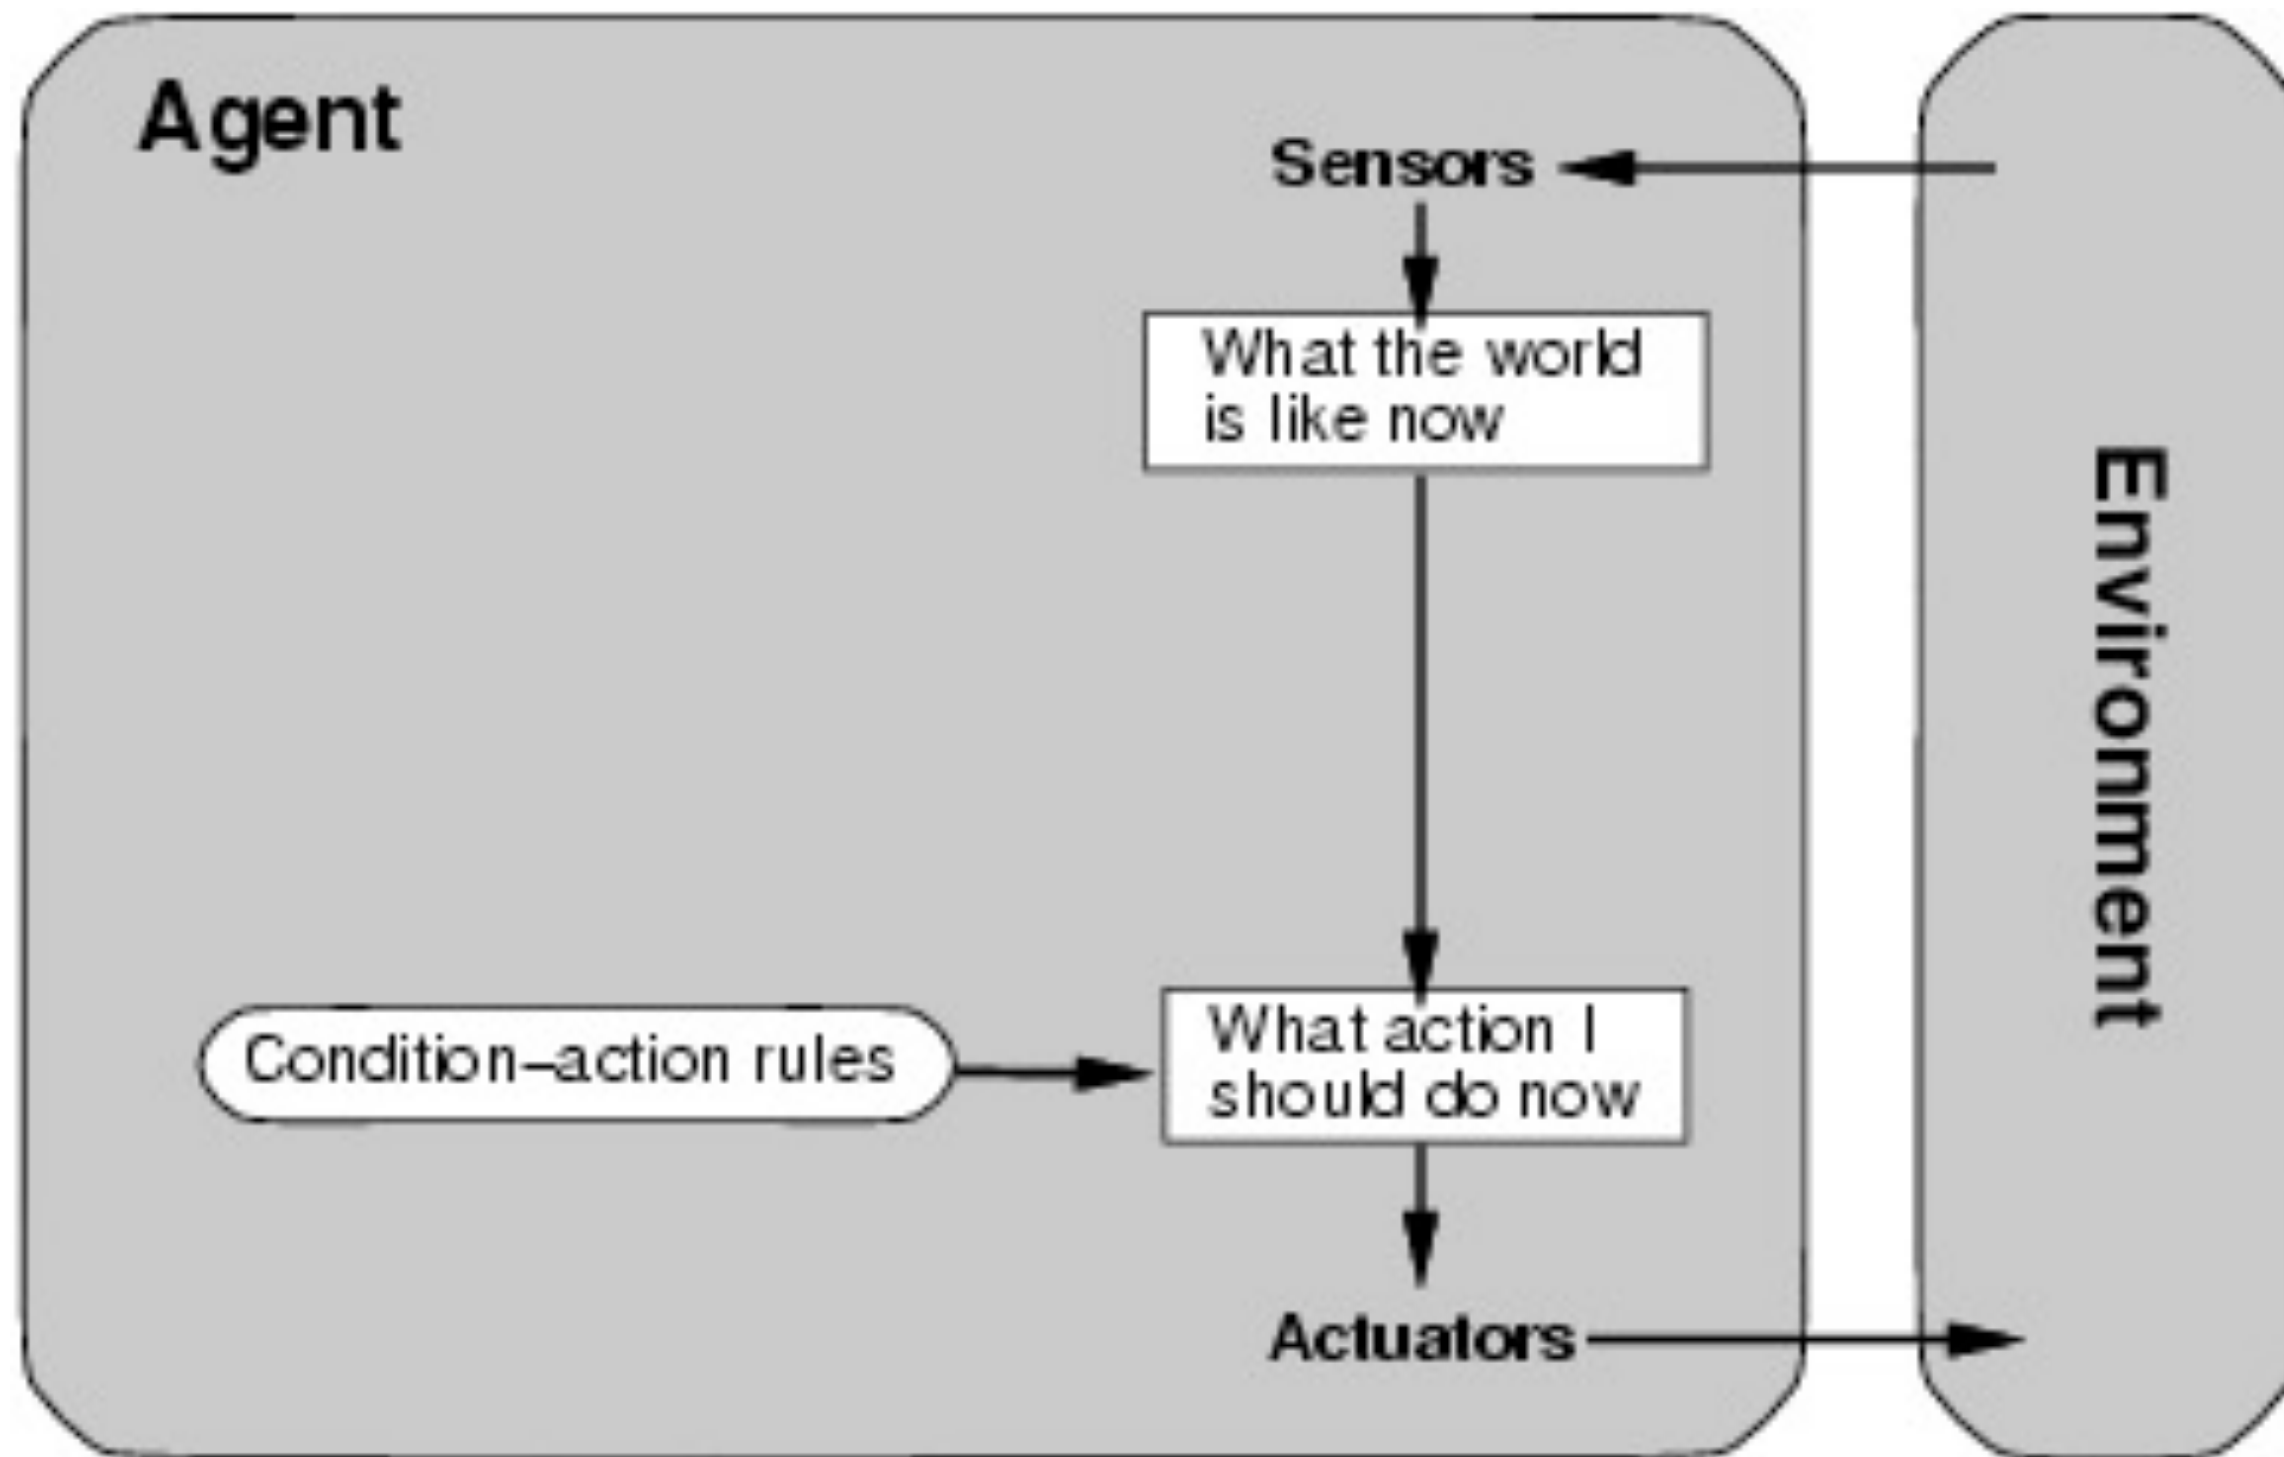
if *status* = *Dirty* then return *Suck*
else if *location* = *A* then return *Right*
else if *location* = *B* then return *Left*

```
(setq joe (make-agent :name 'joe :body (make-agent-body)
                      :program (make-reflex-vacuum-agent-program)))
```

```
(defun make-reflex-vacuum-agent-program ()
  #'(lambda (percept)
      (let ((location (first percept)) (status (second percept)))
        (cond ((eq status 'dirty) 'Suck)
              ((eq location 'A) 'Right)
              ((eq location 'B) 'Left))))))
```

- Four basic types in order of increasing generality:
 - Simple reflex agents
 - Model-based reflex agents
 - Goal-based agents
 - Utility-based agents

Simple Reflex Agents



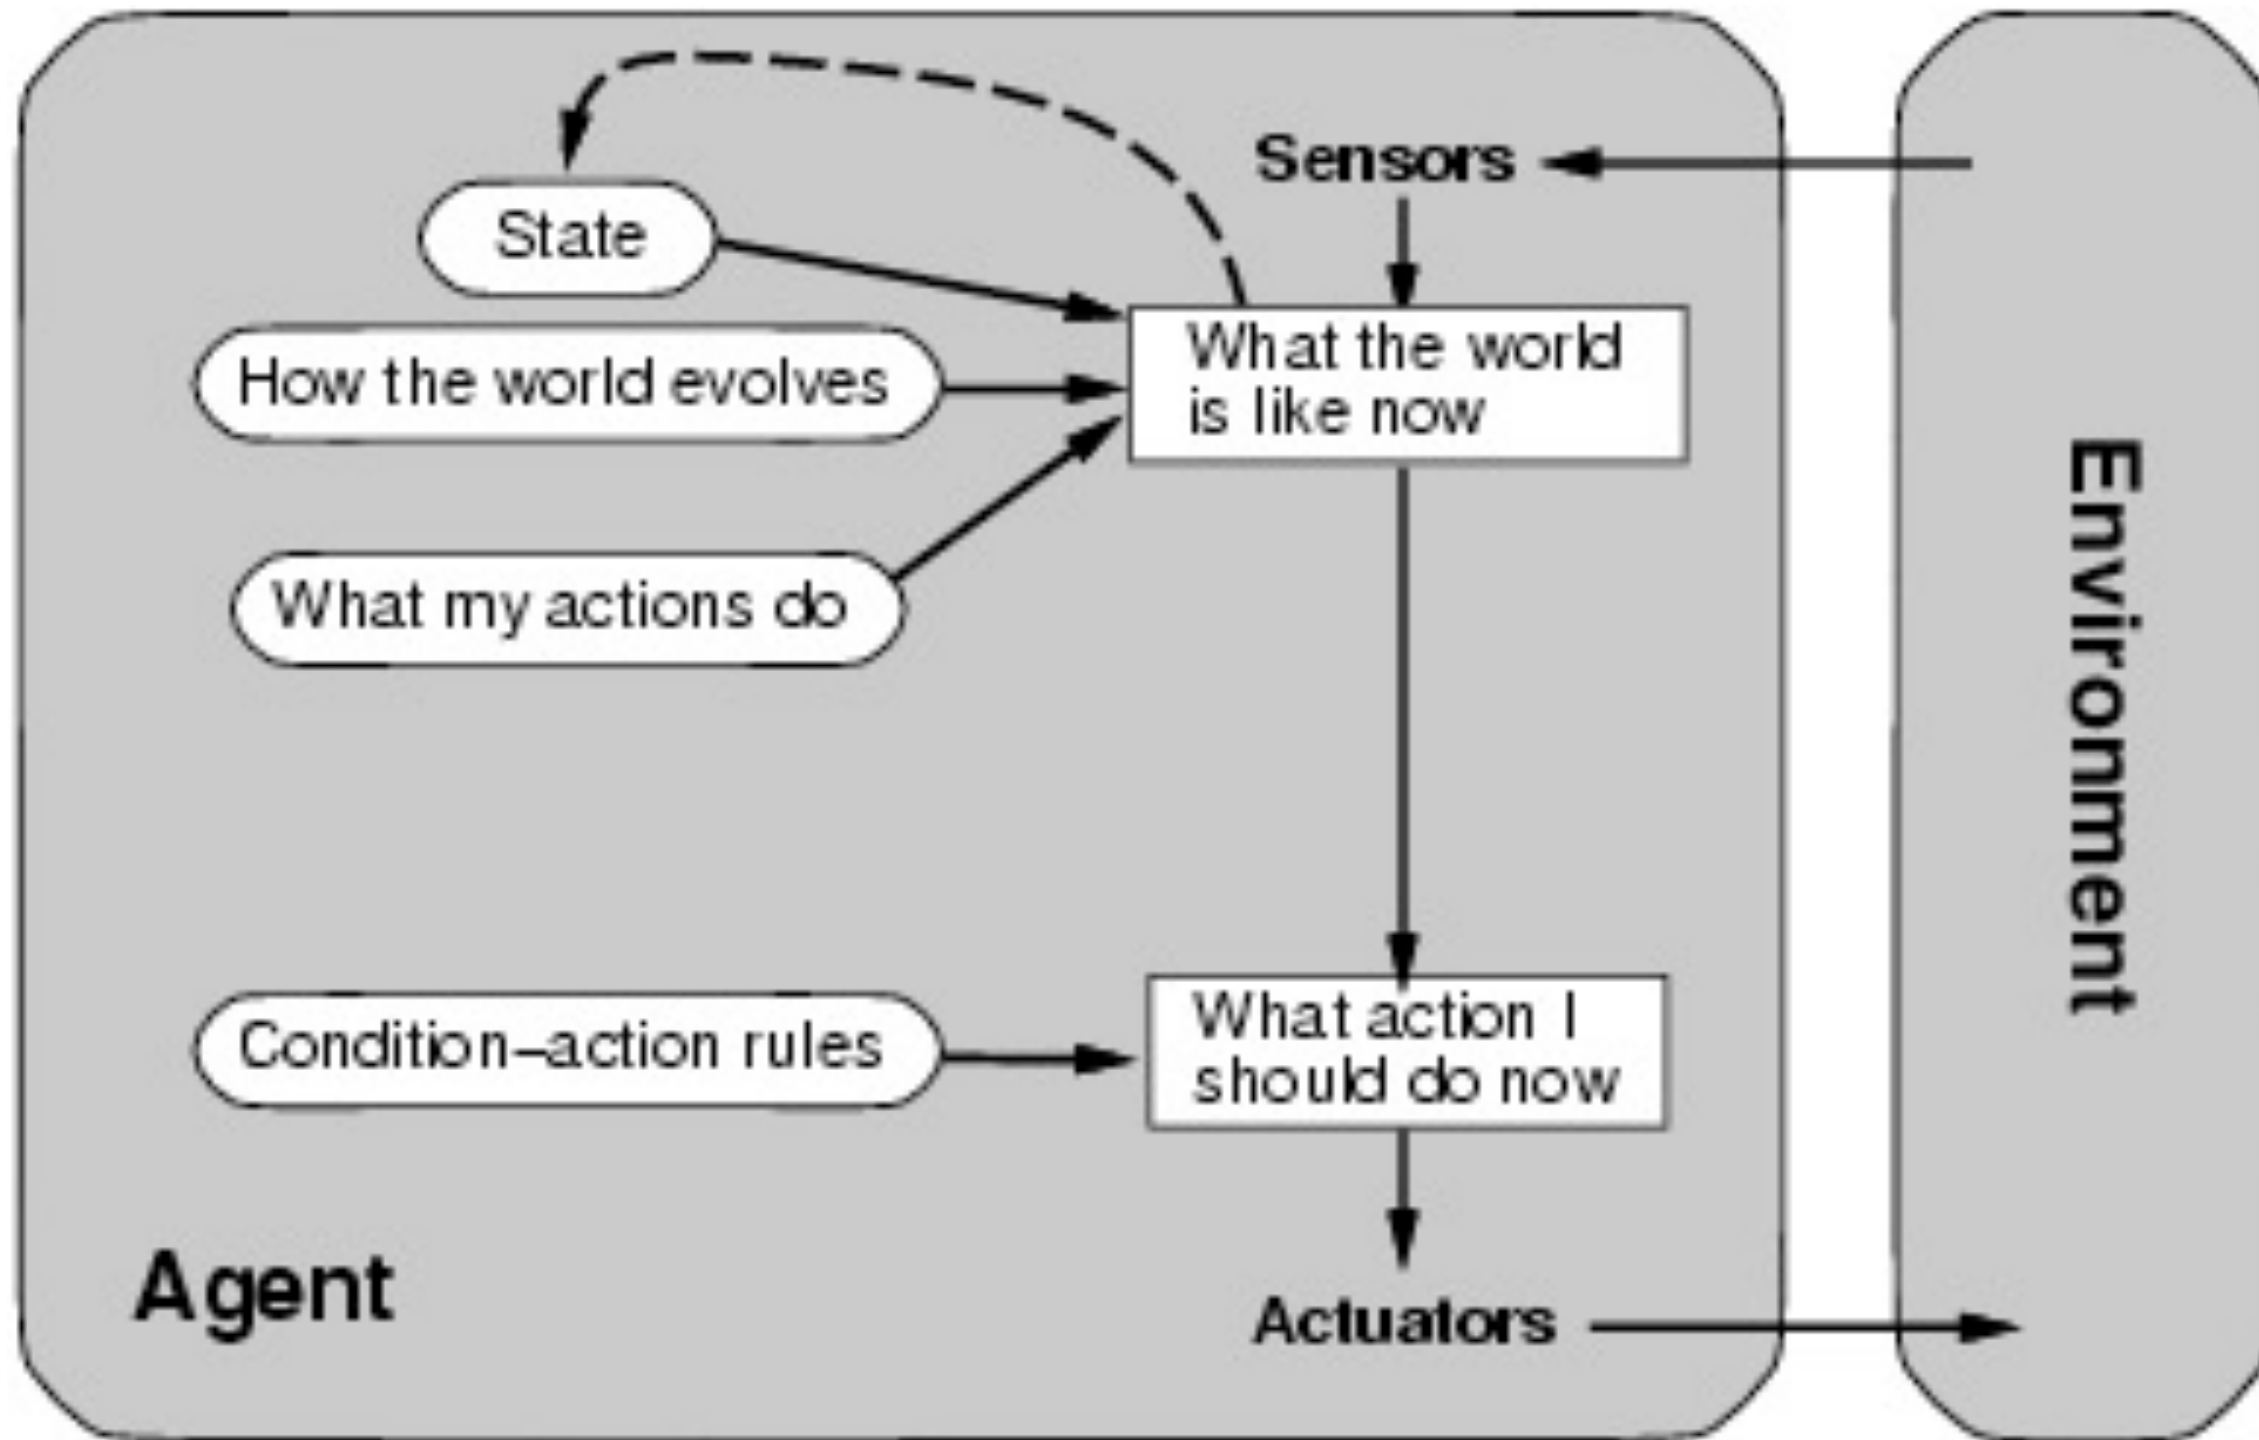
Simple Reflex Agents

```
function REFLEX-VACUUM-AGENT( [location,status]) returns an action
  if status = Dirty then return Suck
  else if location = A then return Right
  else if location = B then return Left
```

```
(setq joe (make-agent :name 'joe :body (make-agent-body)
                      :program (make-reflex-vacuum-agent-program)))

(defun make-reflex-vacuum-agent-program ()
  #'(lambda (percept)
      (let ((location (first percept)) (status (second percept)))
        (cond ((eq status 'dirty) 'Suck)
              ((eq location 'A) 'Right)
              ((eq location 'B) 'Left))))))
```

Model-Based Reflex Agents

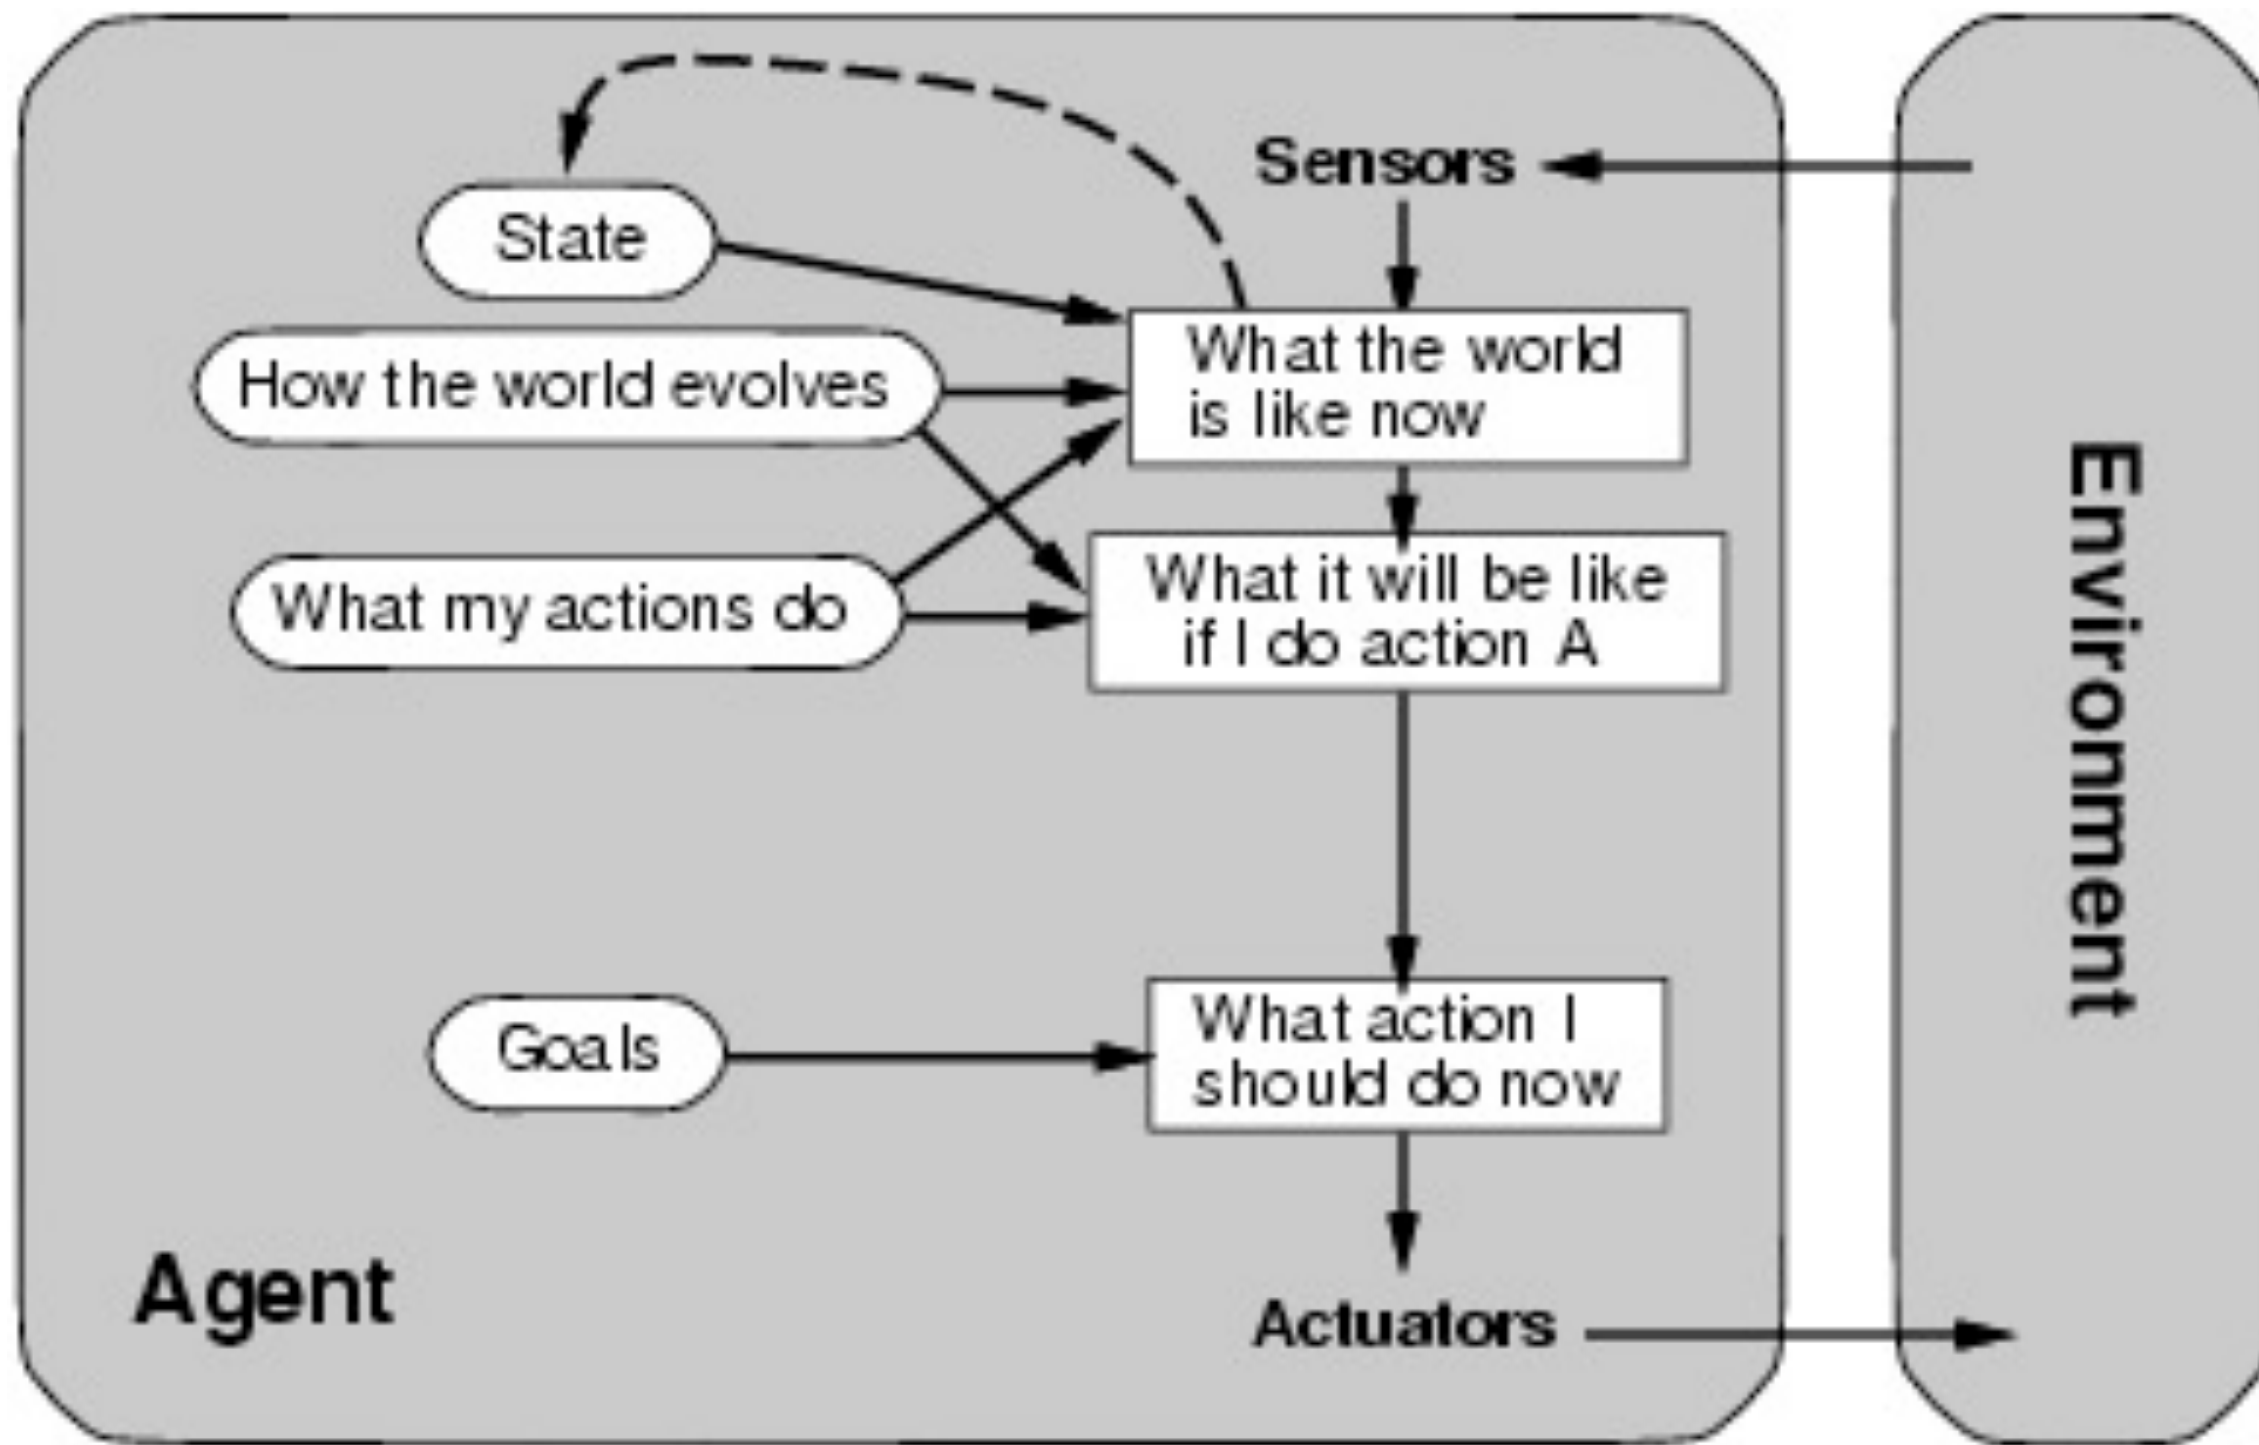


Model-Based Reflex Agents

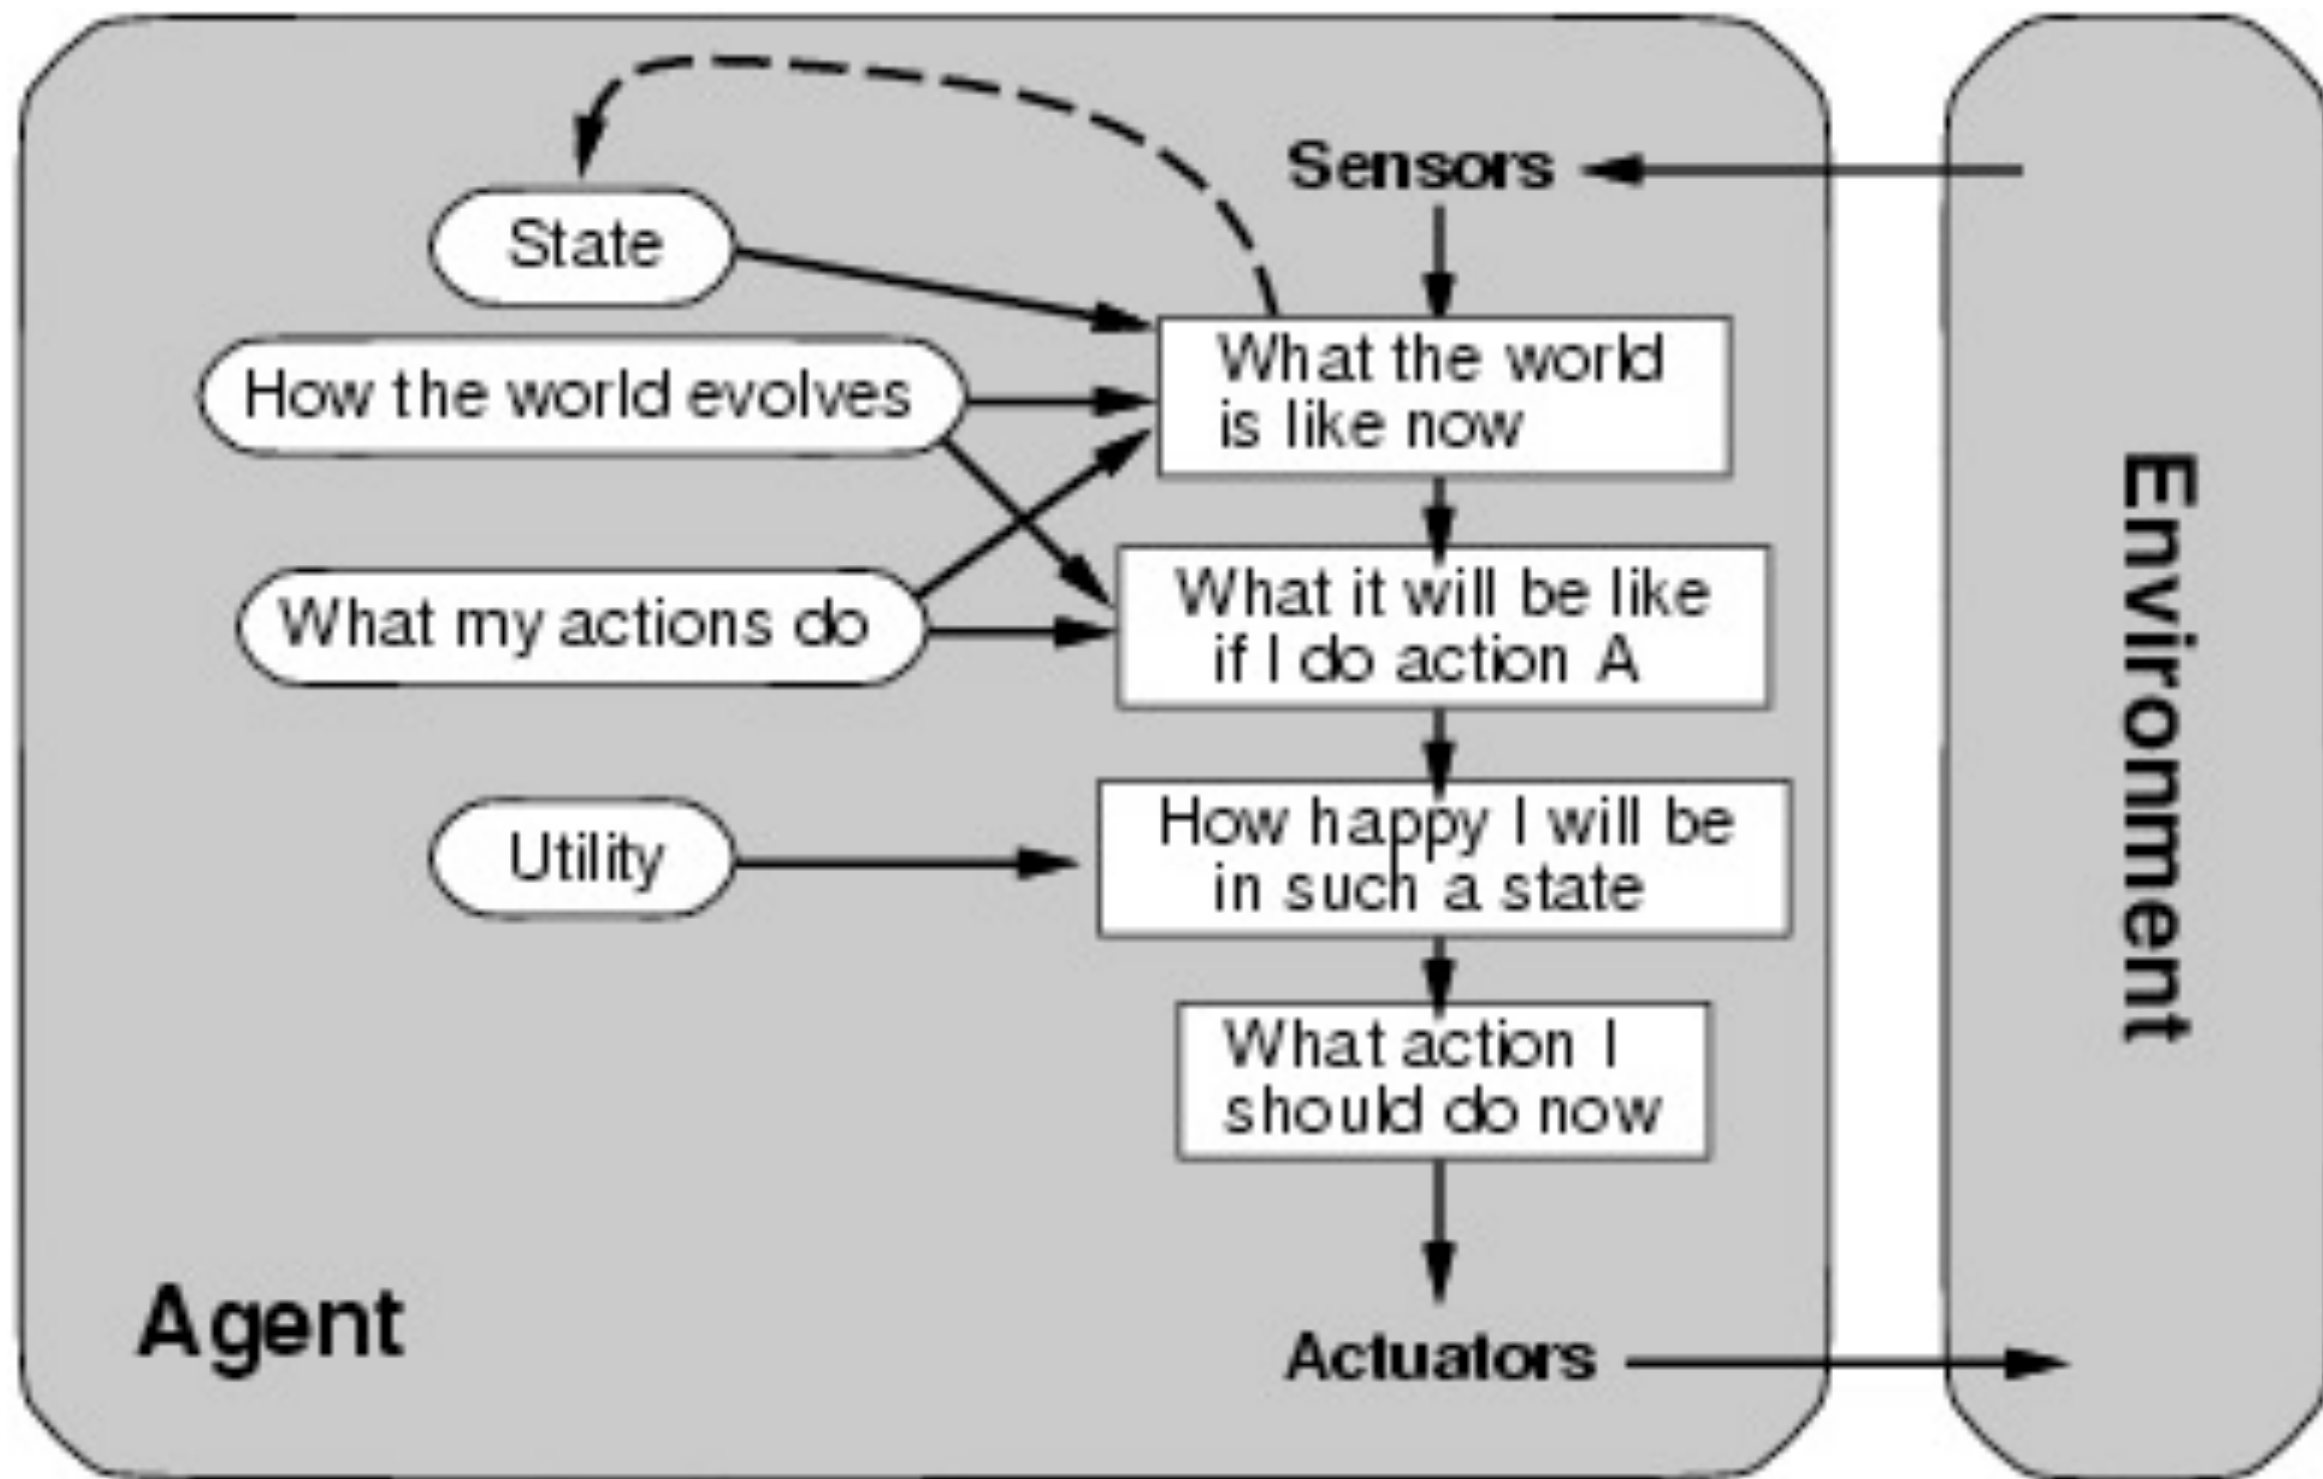
function REFLEX-VACUUM-AGENT(*[location, status]*) returns an action
static: *last_A*, *last_B*, numbers, initially ∞
if *status* = *Dirty* then ...

```
(defun make-reflex-vacuum-agent-with-state-program ()  
  (let ((last-A infinity) (last-B infinity))  
    #'(lambda (percept)  
      (let ((location (first percept)) (status (second percept)))  
        (incf last-A) (incf last-B)  
        (cond  
          ((eq status 'dirty)  
           (if (eq location 'A) (setq last-A 0) (setq last-B 0))  
           'Suck)  
          ((eq location 'A) (if (> last-B 3) 'Right 'NoOp))  
          ((eq location 'B) (if (> last-A 3) 'Left 'NoOp))))))
```


Goal-Based Agents



Utility-Based Agents



- These basic agent types still remain simple and usually cannot deal with more complex environments (dynamic, partially-observable, etc.)
- How do we account e.g. for learning?
 - (see next slide)
 - (see project XPERO)
- Interesting architectures of agents have been investigated in robotics
 - 3T architecture
 - Dynamo, WayFinder
- The architecture of complex agents is still an open research issue!
 - (see projects like DESIRE, ROSTA, XPERO, BRICS)

Learning Agents

