

PLANNING AND SCHEDULING: SEARCH-BASED AGENTS AND BRUTE-FORCE SEARCH

Prof. Dr.-Ing. Gerhard K. Kraetzschmar



Hochschule
Bonn-Rhein-Sieg



Bonn-Aachen
International Center for
Information Technology



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Acknowledgements

- These slides refer to Chapter 3 of the textbook:
S. Russell and P. Norvig:
Artificial Intelligence: A Modern Approach
Prentice Hall, 2003, 2nd Edition (or more recent edition)
- These slides are an adaptation of slides by Min-Yen Kan
- The contributions of these authors are gratefully acknowledged.



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Outline

- Problem-solving agents
- Problem types
- Problem formulation
- Example problems
- Basic search algorithms



Problem-Solving Agents

```
function SIMPLE-PROBLEM-SOLVING-AGENT(percept) returns an action
  static: seq, an action sequence, initially empty
           state, some description of the current world state
           goal, a goal, initially null
           problem, a problem formulation

  state ← UPDATE-STATE(state, percept)
  if seq is empty then do
    goal ← FORMULATE-GOAL(state)
    problem ← FORMULATE-PROBLEM(state, goal)
    seq ← SEARCH(problem)
  action ← FIRST(seq)
  seq ← REST(seq)
  return action
```



Example: Romania

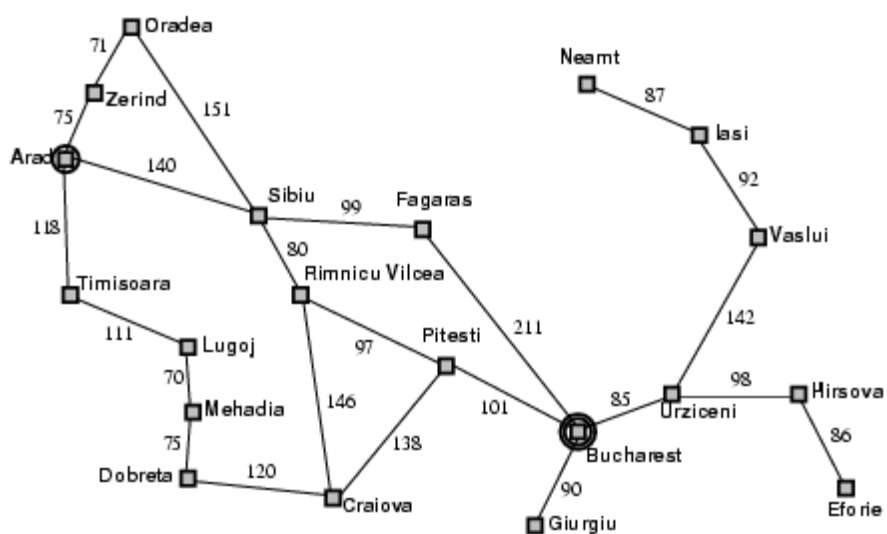
- On holiday in Romania; currently in Arad.
- Flight leaves tomorrow from Bucharest
- Formulate goal:
 - be in Bucharest
- Formulate problem:
 - states: various cities
 - actions: drive between cities
- Find solution:
 - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Example: Romania



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Problem Types

- Deterministic, fully observable
 - → **single-state problem**
 - Agent knows exactly which state it will be in; solution is a sequence
- Non-observable
 - → **sensorless problem** (conformant problem)
 - Agent may have no idea where it is; solution is a sequence
- Nondeterministic and/or partially observable
 - → **contingency problem**
 - percepts provide new information about current state
 - often interleaved search and execution
- Unknown state space
 - → **exploration problem**




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Example: Vacuum World

- Single-state, start in #5.
- Solution?




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Example: Vacuum World

- Single-state, start in #5.
- Solution? **[Right, Suck]**
- Sensorless, start in {1,2,3,4,5,6,7,8} e.g., right goes to {2,4,6,8}
- Solution?




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Example: Vacuum World

- Sensorless, start in {1,2,3,4,5,6,7,8} e.g., right goes to {2,4,6,8}
- Solution? **[Right,Suck,Left,Suck]**
- Contingency
 - Nondeterministic: Suck may dirty a clean carpet
 - Partially observable: location, dirt at current location.
 - Percept: [L, Clean], i.e., start in #5 or #7
- Solution?



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Example: Vacuum World


- Sensorless, start in {1,2,3,4,5,6,7,8}
e.g., right goes to {2,4,6,8}

- Solution? [Right,Suck,Left,Suck]

- Contingency

- Nondeterministic:
Suck may dirty a clean carpet
- Partially observable:
location, dirt at current location.
- Percept:
[L, Clean], i.e., start in #5 or #7

- Solution? [Right, if dirt then Suck]



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Single-State Problem Formulation

- A **problem** is defined by four items:
 - **initial state** e.g., "at Arad"
 - **actions** or **successor function** $S(x)$ = set of action–state pairs
 - e.g., $S(\text{Arad}) = \{ \langle \text{Arad} \rightarrow \text{Zerind}, \text{Zerind} \rangle, \dots \}$
 - **goal test**, can be
 - explicit, e.g., $x = \text{"at Bucharest"}$
 - implicit, e.g., $\text{Checkmate}(x)$
 - **path cost** (additive)
 - e.g., sum of distances, number of actions executed, etc.
 - $c(x,a,y)$ is the step cost, assumed to be ≥ 0
- A **solution** is a sequence of actions leading from the initial state to a goal state



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Selecting a State Space


- Real world is absurdly complex
 - → state space must be **abstracted** for problem solving
- (Abstract) state = set of real states
- (Abstract) action = complex combination of real actions
 - e.g., "Arad → Zerind" represents a complex set of possible routes, detours, rest stops, etc.
- For guaranteed realizability, **any** real state "in Arad" must get to **some** real state "in Zerind"
- (Abstract) solution =
 - set of real paths that are solutions in the real world
- Each abstract action should be "easier" than the original problem



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Vacuum World State Space Graph




- states?
- actions?
- goal test?
- path cost?



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Vacuum World State Space Graph



- states? integer dirt and robot location
- actions? Left, Right, Suck
- goal test? no dirt at all locations
- path cost? 1 per action



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Example: The 8-Puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- states?
- actions?
- goal test?
- path cost?



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Example: The 8-Puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State


- states? locations of tiles
- actions? move blank left, right, up, down
- goal test? = goal state (given)
- path cost? 1 per move
- [Note: optimal solution of n-Puzzle family is **NP-hard**]



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Example: Robotic Assembly



- states? real-valued coordinates of robot joint angles, parts of the object to be assembled
- actions? continuous motions of robot joints
- goal test? complete assembly
- path cost? time to execute



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar


Tree Search Algorithms

- Basic idea:
 - offline, simulated exploration of state space by generating successors of already-explored states (a.k.a. **expanding** states)


```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```



Tree Search Example




Tree Search Example



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Tree Search Example



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Implementation: General Tree Search

```


function TREE-SEARCH( problem, fringe ) returns a solution, or failure
  fringe ← INSERT( MAKE-NODE( INITIAL-STATE[problem] ), fringe )
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST[problem]( STATE[node] ) then return SOLUTION(node)
    fringe ← INSERTALL( EXPAND(node, problem), fringe )

function EXPAND( node, problem ) returns a set of nodes
  successors ← the empty set
  for each action, result in SUCCESSOR-FN[problem]( STATE[node] ) do
    s ← a new NODE
    PARENT-NODE[s] ← node; ACTION[s] ← action; STATE[s] ← result
    PATH-COST[s] ← PATH-COST[node] + STEP-COST(node, action, s)
    DEPTH[s] ← DEPTH[node] + 1
    add s to successors
  return successors
  
```



Implementation: States vs. Nodes

- A **state** is a (representation of) a physical configuration
- A **node** is a data structure constituting part of a search tree
It includes **state**, **parent node**, **action**, **path cost** $g(x)$, **depth**



- The **Expand** function creates new nodes, filling in the various fields and using the **SuccessorFn** of the problem to create the corresponding states.



Search Strategies

- A search strategy is defined by picking the **order of node expansion**
- Different search strategies result from different ways of handling **fringe**
- Strategies are evaluated along the following dimensions:
 - **completeness**: does it always find a solution if one exists?
 - **time complexity**: number of nodes generated
 - **space complexity**: maximum number of nodes in memory
 - **optimality**: does it always find a least-cost solution?
- Time and space complexity are measured in terms of
 - b : maximum branching factor of the search tree
 - d : depth of the least-cost solution
 - m : maximum depth of the state space (may be ∞)




Uninformed Search Strategies

- **Uninformed** search strategies use only the information available in the problem definition
 - Breadth-first search
 - Uniform-cost search
 - Depth-first search
 - Depth-limited search
 - Iterative deepening search



Breadth-First Search

- Principle:
 - Expand shallowest unexpanded node
- Implementation:
 - **fringe** is a FIFO queue, i.e., new successors go at end




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Breadth-First Search

- Principle:
 - Expand shallowest unexpanded node
- Implementation:
 - **fringe** is a FIFO queue, i.e., new successors go at end




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Breadth-First Search

- Principle:
 - Expand shallowest unexpanded node
- Implementation:
 - **fringe** is a FIFO queue, i.e., new successors go at end




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Breadth-First Search

- Principle:
 - Expand shallowest unexpanded node
- Implementation:
 - **fringe** is a FIFO queue, i.e., new successors go at end



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Properties of Breadth-First Search

- Complete? Yes (if b is finite)
- Time? $1+b+b^2+b^3+\dots +b^d + b(b^{d+1}-b) = O(b^{d+1})$
- Space? $O(b^{d+1})$ (keeps every node in memory)
- Optimal? Yes (if cost = 1 per step)
- **Space** is the bigger problem (more than time)



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Uniform-Cost Search

- Principle:
 - Expand least-cost unexpanded node
- Implementation:
 - **fringe** = queue ordered by path cost
- Equivalent to breadth-first if step costs all equal
- Complete? Yes, if step cost $\geq \epsilon$
- Time? # of nodes with $g \leq$ cost of optimal solution,
 $O(b^{\text{ceiling}(C^*/\epsilon)})$
where C^* is the cost of the optimal solution
- Space? # of nodes with $g \leq$ cost of optimal solution,
 $O(b^{\text{ceiling}(C^*/\epsilon)})$
- Optimal? Yes – nodes expanded in increasing order of $g(n)$




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Depth-First Search

- Principle:
 - Expand deepest unexpanded node
- Implementation:
 - **fringe** = LIFO queue, i.e. put successors at front



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Properties of Depth-First Search

- Complete? No: fails in infinite-depth spaces, spaces with loops
 - Modify to avoid repeated states along path
 - → complete in finite spaces
- Time? $O(b^m)$: terrible if m is much larger than d
 - but if solutions are dense, may be much faster than breadth-first
- Space? $O(bm)$, i.e., linear space!
- Optimal? No



Depth-Limited Search

- Principle:
 - depth-first search with depth limit l , i.e., nodes at depth l have no successors
- Recursive implementation:

```
function DEPTH-LIMITED-SEARCH(problem, limit) returns soln/fail/cutoff
  RECURSIVE-DLS(MAKE-NODE(INITIAL-STATE[problem]), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns soln/fail/cutoff
  cutoff-occurred? ← false
  if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
  else if DEPTH[node] = limit then return cutoff
  else for each successor in EXPAND(node, problem) do
    result ← RECURSIVE-DLS(successor, problem, limit)
    if result = cutoff then cutoff-occurred? ← true
    else if result ≠ failure then return result
  if cutoff-occurred? then return cutoff else return failure
```



Iterative Deepening Search

- Principle:
 - Perform depth-limited search with increasing limit l
- Implementation:

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or fail-  
ure  
  inputs: problem, a problem  
  for depth  $\leftarrow 0$  to  $\infty$  do  
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)  
    if result  $\neq$  cutoff then return result
```



Iterative Deepening Search $l=0$

Limit = 0



Iterative Deepening Search $l = 1$

Limit = 1




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Iterative Deepening Search $l = 2$


Limit = 2



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Iterative Deepening Search I = 3



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Iterative Deepening Search

- Number of nodes generated in a depth-limited search to depth d with branching factor b :

$$N_{\text{DLS}} = b^0 + b^1 + b^2 + \dots + b^{d-2} + b^{d-1} + b^d$$

- Number of nodes generated in an iterative deepening search to depth d with branching factor b :

$$N_{\text{IDS}} = (d+1)b^0 + d b^1 + (d-1)b^2 + \dots + 3b^{d-2} + 2b^{d-1} + 1b^d$$

- Example: for $b = 10$, $d = 5$,

$$N_{\text{DLS}} = 1 + 10 + 100 + 1,000 + 10,000 + 100,000 = 111,111$$

$$N_{\text{IDS}} = 6 + 50 + 400 + 3,000 + 20,000 + 100,000 = 123,456$$

- Overhead = $(123,456 - 111,111)/111,111 = 11\%$



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Properties of Iterative Deepening Search

- Complete? Yes
- Time? $(d+1)b^0 + d b^1 + (d-1)b^2 + \dots + b^d = O(b^d)$
- Space? $O(bd)$
- Optimal? Yes, if step cost = 1



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Summary of Algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening
Complete?	Yes	Yes	No	No	Yes
Time	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$
Space	$O(b^{d+1})$	$O(b^{\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$
Optimal?	Yes	Yes	No	No	Yes




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Repeated States

- Failure to detect repeated states can turn a linear problem into an exponential one!



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Graph Search

```

function GRAPH-SEARCH(problem, fringe) returns a solution, or failure
    closed ← an empty set
    fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
    loop do
        if fringe is empty then return failure
        node ← REMOVE-FRONT(fringe)
        if GOAL-TEST[problem](STATE[node]) then return SOLUTION(node)
        if STATE[node] is not in closed then
            add STATE[node] to closed
            fringe ← INSERTALL(EXPAND(node, problem), fringe)
    
```



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Summary

- Problem formulation usually requires abstracting away real-world details to define a state space that can feasibly be explored
- Variety of uninformed search strategies
- Iterative deepening search uses only linear space and not much more time than other uninformed algorithms

