

PLANNING AND SCHEDULING: INFORMED SEARCH

Prof. Dr.-Ing. Gerhard K. Kraetzschmar



Hochschule
Bonn-Rhein-Sieg



Bonn-Aachen
International Center for
Information Technology



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Acknowledgements

- These slides refer to Chapter 4 of the textbook:
S. Russell and P. Norvig:
Artificial Intelligence: A Modern Approach
Prentice Hall, 2003, 2nd Edition (or more recent edition)
- These slides are an adaptation of slides by Min-Yen Kan
- The contributions of these authors are gratefully acknowledged.



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Outline

- Best-first search
- Greedy best-first search
- A* search
- Heuristics
- Local search algorithms
- Hill-climbing search
- Simulated annealing search
- Local beam search
- Genetic algorithms



Review: Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
```

- A search strategy is defined by picking the order of node expansion



Best-First Search

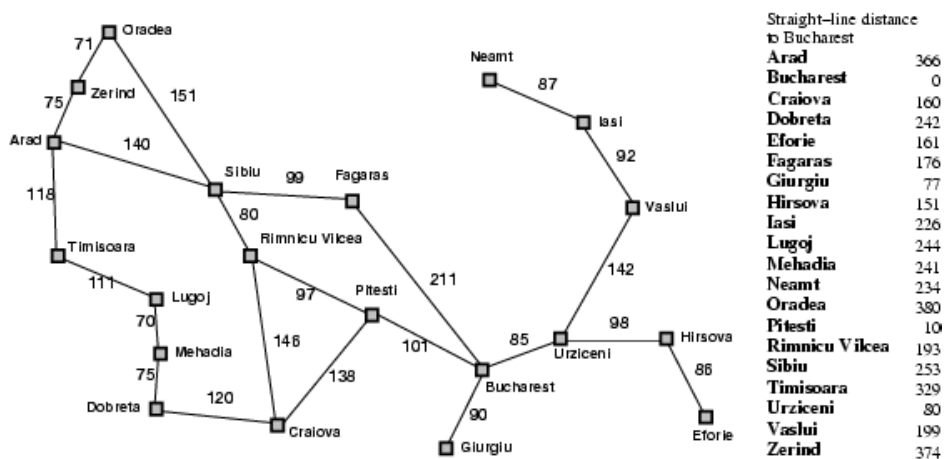
- Idea: use an **evaluation function** $f(n)$ for each node
 - Estimate of "desirability"
 - Expand most desirable unexpanded node
- Implementation:
 - Order the nodes in fringe in decreasing order of desirability
- Special cases:
 - Greedy Best-First Search
 - A* Search



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Romania With Step Costs In km



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Greedy Best-First Search

- Evaluation function $f(n) = h(n)$ (heuristic)
 - = estimate of cost from n to goal
- For example:
 - $h_{SLD}(n)$ = straight-line distance from n to Bucharest
- Greedy best-first search expands the node that **appears** to be closest to goal



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Greedy Best-First Search Example



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar


Greedy Best-First Search Example



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar


Greedy Best-First Search Example



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Greedy Best-First Search Example



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Properties of Greedy Best-First Search

- Complete? No
 - can get stuck in loops, e.g., Iasi → Neamt → Iasi → Neamt → ...
- Time? $O(b^m)$
 - but a good heuristic can give dramatic improvement
- Space? $O(b^m)$
 - keeps all nodes in memory
- Optimal? No



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

A* Search

- Idea:
 - Avoid expanding paths that are already expensive
- Evaluation function $f(n) = g(n) + h(n)$
 - $g(n)$ = cost so far to reach n
 - $h(n)$ = estimated cost from n to goal
 - $f(n)$ = estimated total cost of path through n to goal



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

A* Search Example

▶ Arad
366=0+366



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar


A* Search Example



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar


A* Search Example



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar


A* Search Example



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

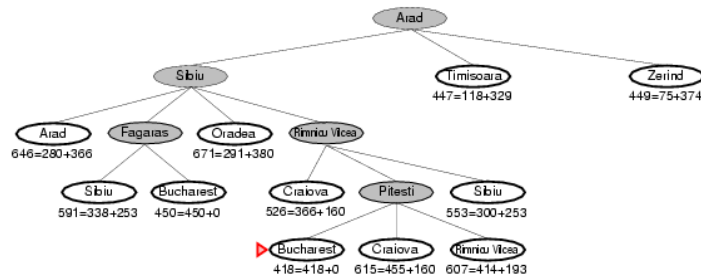
A* Search Example



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

A* Search Example



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Admissible Heuristics

- A heuristic $h(n)$ is **admissible** if for every node n , $h(n) \leq h^*(n)$, where $h^*(n)$ is the **true** cost to reach the goal state from n .
- An admissible heuristic **never overestimates** the cost to reach the goal, i.e., it is **optimistic**.
- Example: $h_{SLD}(n)$ (SLD = Straight-Line Distance)
(never overestimates the actual road distance)
- **Theorem**: If $h(n)$ is admissible, A* using **TREE-SEARCH** is optimal




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Optimality of A* (Proof)

- Suppose some suboptimal goal G_2 has been generated and is in the fringe. Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G .



- $f(G_2) = g(G_2)$ since $h(G_2) = 0$
- $g(G_2) > g(G)$ since G_2 is suboptimal
- $f(G) = g(G)$ since $h(G) = 0$
- $f(G_2) > f(G)$ from above




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Optimality of A* (Proof)

- Suppose some suboptimal goal G_2 has been generated and is in the fringe. Let n be an unexpanded node in the fringe such that n is on a shortest path to an optimal goal G .



- $f(G_2) > f(G)$ from above
- $h(n) \leq h^*(n)$ since h is admissible
- $g(n) + h(n) \leq g(n) + h^*(n)$
- $f(n) \leq f(G)$
- Hence $f(G_2) > f(n)$, and A* will never select G_2 for expansion



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Consistent Heuristics


- A heuristic is **consistent** if for every node n , every successor n' of n generated by any action a ,

$$h(n) \leq c(n,a,n') + h(n')$$

- If h is consistent, we have

$$\begin{aligned} f(n') &= g(n') + h(n') \\ &= g(n) + c(n,a,n') + h(n') \\ &\geq g(n) + h(n) \\ &= f(n) \end{aligned}$$

i.e., $f(n)$ is non-decreasing along any path.



- Theorem:** If $h(n)$ is consistent, A* using **GRAPH-SEARCH** is optimal




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Optimality of A*

- A* expands nodes in order of increasing f value
- Gradually adds "f-contours" of nodes
- Contour i has all nodes with $f = f_i$, where $f_i < f_{i+1}$



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Properties of A*

- Complete? Yes
 - (unless there are infinitely many nodes with $f \leq f(G)$)
- Time? Exponential
- Space? Keeps all nodes in memory
- Optimal? Yes



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Admissible Heuristics

E.g., for the 8-puzzle:

- $h1(n)$ = number of misplaced tiles
- $h2(n)$ = total Manhattan distance

(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h1(S) = ?$
- $h2(S) = ?$



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Admissible Heuristics

E.g., for the 8-puzzle:

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total Manhattan distance

(i.e., no. of squares from desired location of each tile)

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State

- $h_1(S) = ? 8$
- $h_2(S) = ? 3+1+2+2+2+3+3+2 = 18$



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Dominance

- If $h_2(n) \geq h_1(n)$ for all n (both admissible), then h_2 **dominates** h_1
- If h_2 dominates h_1 , h_2 is better for search
- Typical search costs (average number of nodes expanded):
 - $d=12$ IDS = 3,644,035 nodes
 $A^*(h_1) = 227$ nodes
 $A^*(h_2) = 73$ nodes
 - $d=24$ IDS = too many nodes
 $A^*(h_1) = 39,135$ nodes
 $A^*(h_2) = 1,641$ nodes



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Relaxed Problems

- A problem with fewer restrictions on the actions is called a **relaxed problem**
- **The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem**
- If the rules of the 8-puzzle are relaxed so that a tile can move **anywhere**, then $h_1(n)$ gives the shortest solution
- If the rules are relaxed so that a tile can move to **any adjacent square**, then $h_2(n)$ gives the shortest solution



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Local Search Algorithms

- In many optimization problems, the **path** to the goal is irrelevant;
the goal state itself is the solution
- State space = set of "complete" configurations
- Find configuration satisfying constraints, e.g., n-queens
- In such cases, we can use **local search algorithms**
- Keep a single "current" state, try to improve it



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Example: n-Queens

- Put n queens on an $n \times n$ board with no two queens on the same row, column, or diagonal



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Hill-Climbing Search

- "Like climbing Mt Everest in thick fog with amnesia"

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                  neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] ≤ VALUE[current] then return STATE[current]
    current ← neighbor
```




Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Hill-Climbing Search

- Problem: depending on the initial state, we can get stuck in local maxima



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Hill-Climbing Search: 8-Queens Problem

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♚	13	16	13	16
♚	14	17	15	♚	14	16	16
17	♚	16	18	15	♚	15	♚
18	14	♚	15	15	14	♚	16
14	14	13	17	12	14	12	18


- h = number of pairs of queens that are attacking each other, either directly or indirectly
- $h = 17$ for the above state



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Hill-Climbing Search: 8-Queens Problem



- A local minimum with $h = 1$



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Simulated Annealing Search

- Idea:
 - Escape local maxima by allowing some "bad" moves
 - But gradually decrease the frequency of "bad" moves

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
         schedule, a mapping from time to "temperature"
  local variables: current, a node
                  next, a node
                  T, a "temperature" controlling prob. of downward steps

  current ← MAKE-NODE(INITIAL-STATE[problem])
  for t ← 1 to ∞ do
    T ← schedule[t]
    if T = 0 then return current
    next ← a randomly selected successor of current
     $\Delta E \leftarrow \text{VALUE}[\textit{next}] - \text{VALUE}[\textit{current}]$ 
    if  $\Delta E > 0$  then current ← next
    else current ← next only with probability  $e^{\Delta E/T}$ 
```



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Properties of Simulated Annealing Search

- One can prove:
 - If T decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1
- Widely used in VLSI layout, airline scheduling, etc.



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Local Beam Search

- Keep track of k states rather than just one
- Start with k randomly generated states
- At each iteration, all the successors of all k states are generated
- If any one is a goal state, stop; else select the k best successors from the complete list and repeat.



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Genetic Algorithms


- A successor state is generated by combining two parent states
- Start with k randomly generated states (a **population**)
- A state is represented as a string over a finite alphabet (often a string of 0s and 1s)
- Evaluation function (**fitness function**).
Higher values for better states.
- Produce the next generation of states by
 - **Selection**,
 - **Crossover**, and
 - **Mutation**



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Genetic Algorithms




- Fitness function:
 - number of non-attacking pairs of queens (min = 0, max = $8 \times 7/2 = 28$)
 - $24/(24+23+20+11) = 31\%$
 - $23/(24+23+20+11) = 29\%$ etc



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Genetic Algorithms



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar