

PLANNING AND SCHEDULING: HIERARCHICAL TASK NETWORK PLANNING

Prof. Dr.-Ing. Gerhard K. Kraetzschmar



Hochschule
Bonn-Rhein-Sieg



Bonn-Aachen
International Center for
Information Technology



Hochschule
Bonn-Rhein-Sieg

© 2009 Gerhard K. Kraetzschmar

Acknowledgements

These slides are based on those slides by Dana Nau, Gerhard Wickler, Hai Hoang, José Luis Ambite

Several improvements were applied by Iman Awaad

Motivation

- We may already have an idea how to go about solving problems in a planning domain
- E.g.: travel to a destination that's far away:
 - Domain-independent planner:
 - many combinations of vehicles and routes
 - Experienced human: small number of “recipes”, e.g. for flying:
 - buy ticket from local airport to remote airport
 - travel to local airport
 - fly to remote airport
 - travel to final destination
- How to enable planning systems to make use of such recipes?

Control Rules v HTN Planning

1. Control rules (Chapter 10):
 - Write rules to prune every action that **does not** fit the recipe
2. Hierarchical Task Network (HTN) planning:
 - Describe the actions and subtasks that **do** fit the recipe

Objective of HTN planning: perform a given set of tasks

- Inputs include:
 - *Operators*: that can directly perform a *primitive* task
 - *Methods*: recipes for decomposing a complex/*non-primitive task* into simpler non-primitive or primitive subtasks
- Planning process:
 - *Decompose* non-primitive tasks recursively until primitive tasks are reached

Hierarchical Decomposition & Problem Reduction

To get to a conference in ?x, get to the airport, take a plane to ?x, then go to the conference hotel

- To get to the airport, either drive or take a cab
- If you have money for the taxi fare:
- Enter the cab, say “I want to go to ?y”, wait until you are at ?y, pay the fare, then exit the taxi”
- Idea is to capture the hierarchical structure of the planning domain
 - contains complex tasks and schemas for reducing them.
- Reduction schemas:
 - given by the designer
 - express preferred ways to accomplish a task

Outline

- Main idea behind HTN planning
- STNs: Representation and planning algorithms
 - Total order
 - Partial order
- Generalizing the formalism and algorithm to HTN
- Expressivity: comparison to classical planning and control rules
- Experimental Results

HTN Planning

- A type of *problem reduction*
- Decompose *tasks* into *subtasks*
- Handle constraints (e.g., taxi not good for long distances)
- Resolve interactions (e.g., take taxi early enough to catch plane)
- If necessary, backtrack and try other decompositions

travel(UMD, LAAS)

get-ticket(BWI, Toulouse)

go to Orbitz
find-flights(BWI, Toulouse)
buy-ticket(BWI, Toulouse)

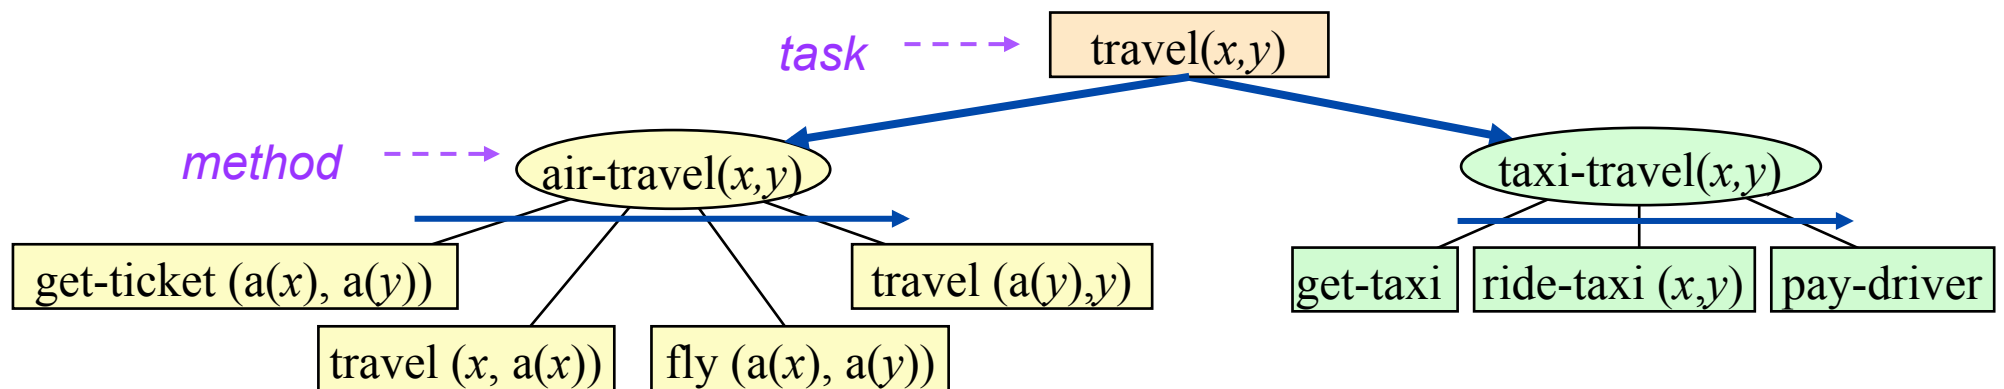
travel(UMD, BWI)

get-taxi
ride-taxi(UMD, BWI)
pay-driver

fly(BWI, Toulouse)

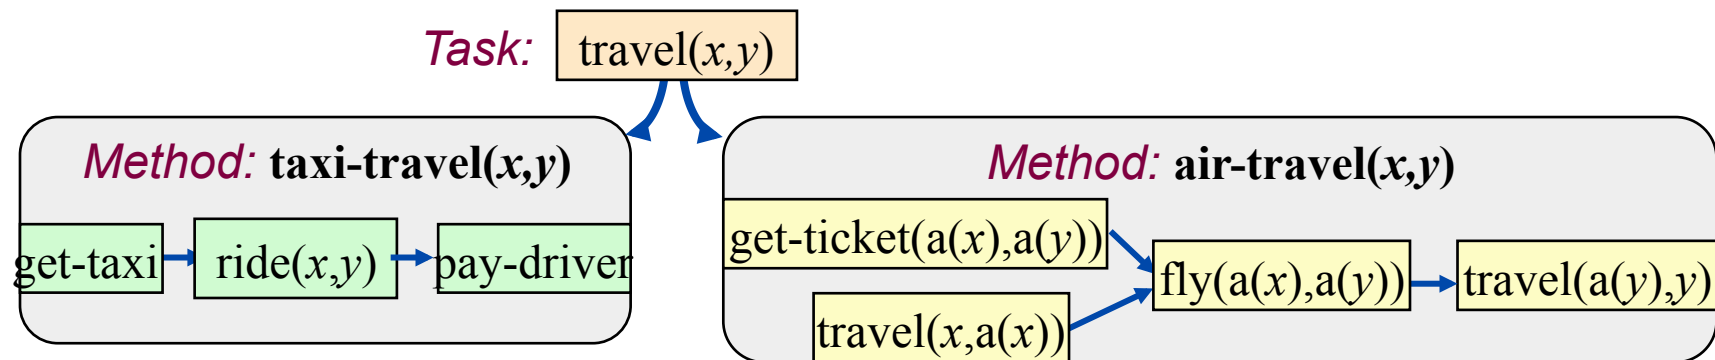
travel(Toulouse, LAAS)

get-taxi
ride-taxi(Toulouse, LAAS)
pay-driver



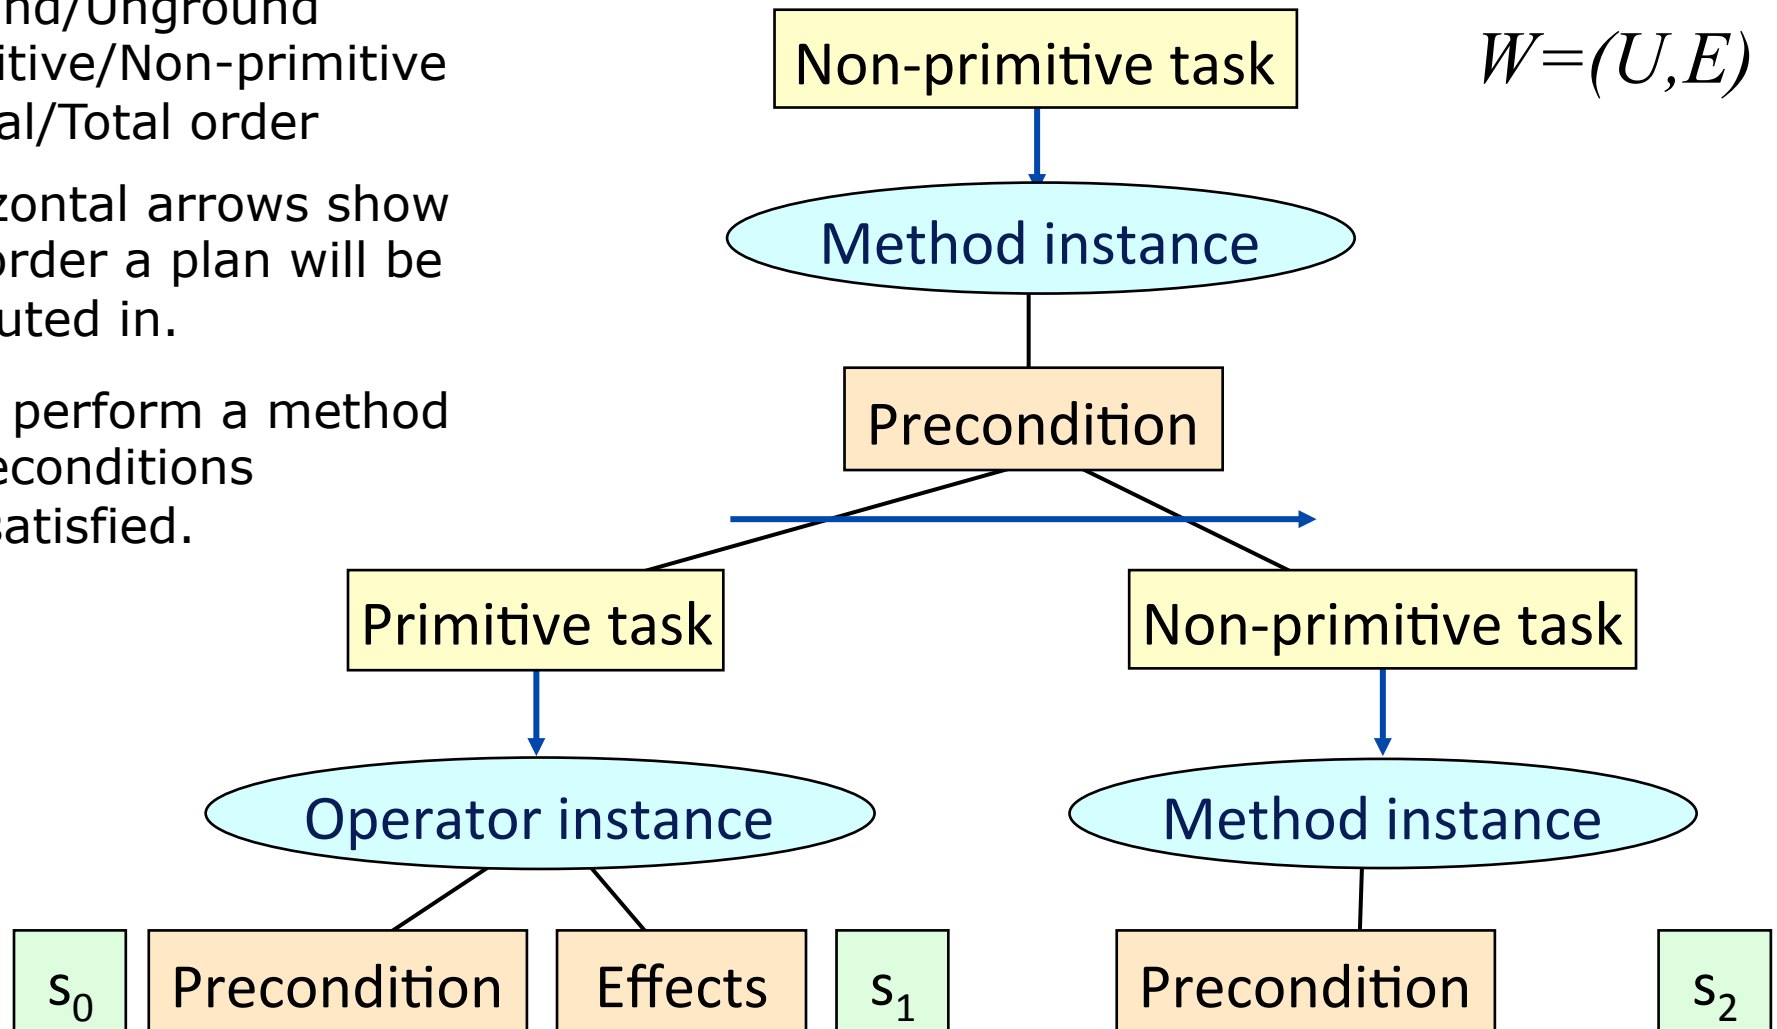
HTN Planning

- HTN planners may be domain-specific
 - e.g., see Chapters 20 (robotics) and 23 (bridge)
- Or they may be domain-configurable
 - Domain-independent planning engine
 - Domain description defining operators and also methods
- Problem description
 - domain description, initial state, initial task network



HTN Planning: Task Networks

- Ground/Unground
- Primitive/Non-primitive
- Partial/Total order
- Horizontal arrows show the order a plan will be executed in.
- Only perform a method if preconditions are satisfied.



HTN v what we've seen so far

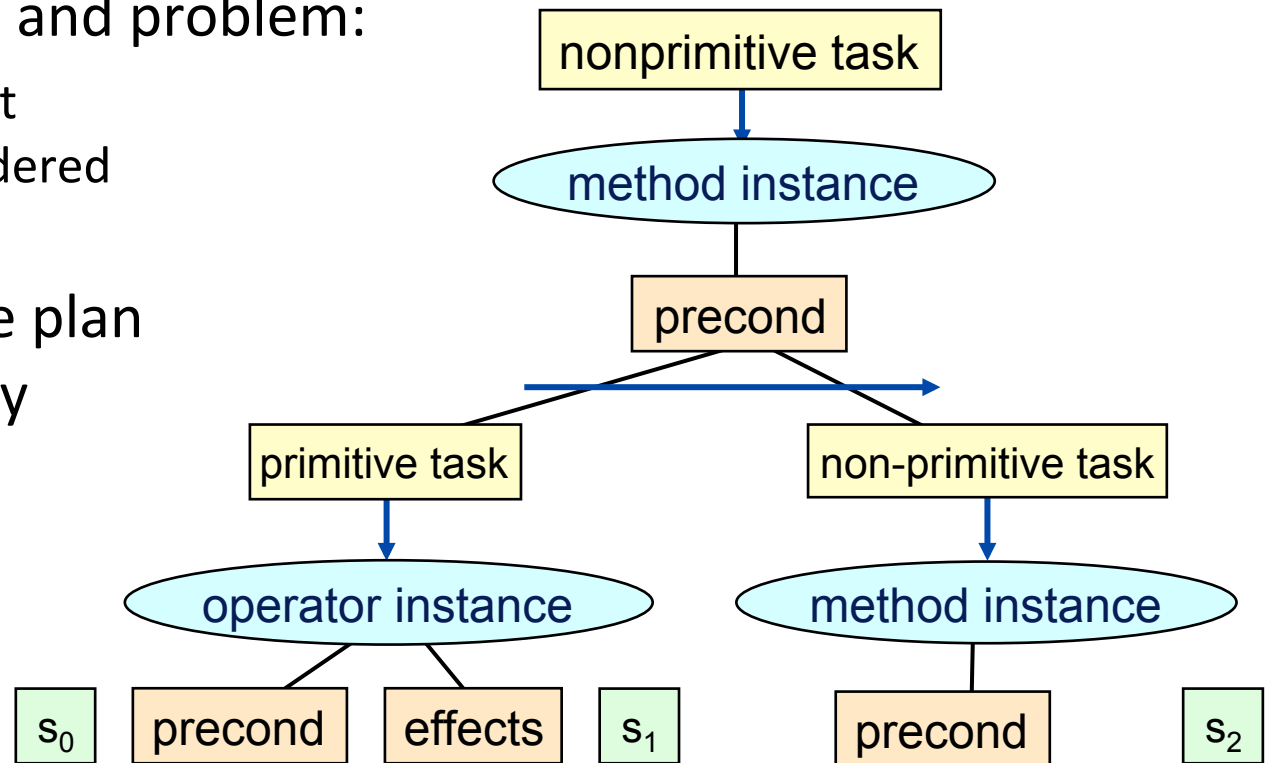
- What stays the same:
 - Each state of the world is represented by a set of atoms
 - Each action corresponds to a deterministic state transition
 - Terms, literals, operators, actions, plans have same meaning
 - E.g. (block b1) (block b2) (block b3) (block b4) (on-table b1) (on b2 b1) (clear b2) (on-table b3) (on b4 b3) (clear b4)
- What's new:
 - Perform a set of tasks not achieve a set of goals
 - *Methods* describing ways in which tasks can be performed
 - Organized collections of tasks and subtasks called *task networks*

Simple Task Network (STN) Planning

- A special case of HTN planning
- States and operators
 - The same as in classical planning
- *Task*: an expression of the form $t(u_1, \dots, u_n)$
 - t is a *task symbol*, and each u_i is a term
- Two kinds of task symbols (and tasks):
 - *Primitive*: tasks that we know how to execute directly
 - task symbol is an operator name
 - *Non-primitive*: tasks that must be decomposed into subtasks
 - use methods (next slide)

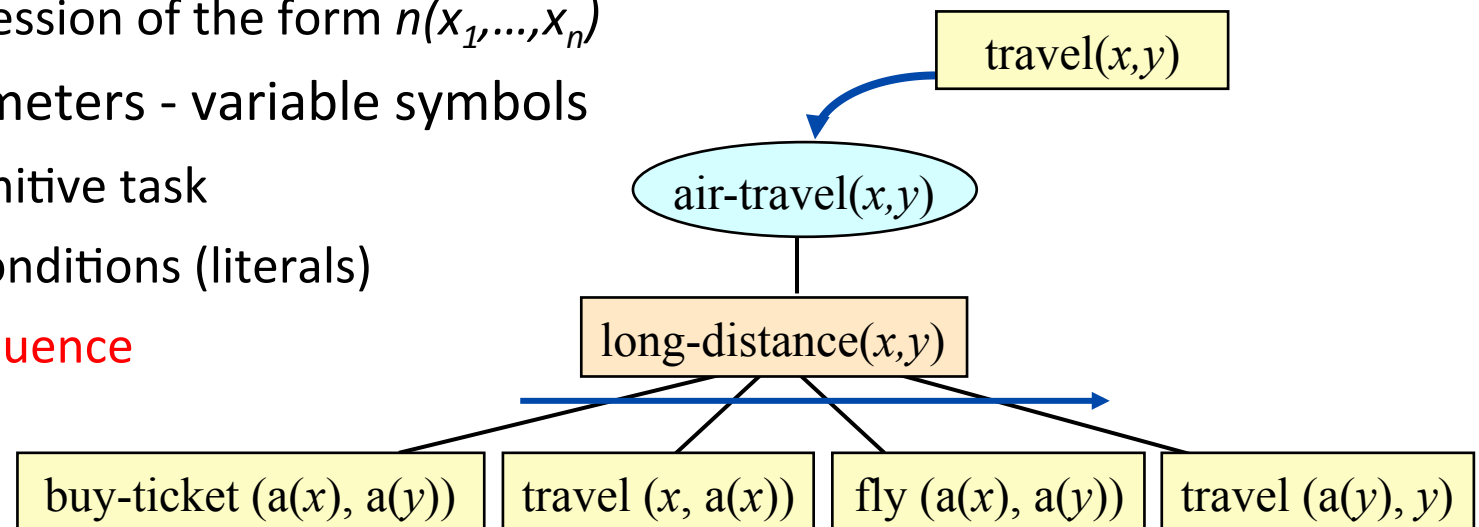
STN: Domains, Planning Problems, Solutions

- Domain: methods, operators: $D=(O,M)$
- Problem: initial state, initial task network, operators, methods:
 $P=(S_0, w_j, O, M)$
- Total-order STN domain and problem:
 - Same as above except that all methods are totally ordered
- Solution: any executable plan that can be generated by recursively applying
 - methods to non-primitive tasks
 - operators to primitive tasks



STN: Methods (Totally-ordered)

- Totally-ordered method: a 4-tuple
 $m = (\text{name}(m), \text{task}(m), \text{precond}(m), \text{subtasks}(m))$
 - $\text{name}(m)$: an expression of the form $n(x_1, \dots, x_n)$
 - x_1, \dots, x_n are parameters - variable symbols
 - $\text{task}(m)$: a nonprimitive task
 - $\text{precond}(m)$: preconditions (literals)
 - $\text{subtasks}(m)$: a **sequence** of tasks $\langle t_1, \dots, t_k \rangle$



`air-travel(x,y)`

task: `travel(x,y)`

precond: `long-distance(x,y)`

subtasks: $\langle \text{buy-ticket}(a(x), a(y)), \text{travel}(x, a(x)), \text{fly}(a(x), a(y)), \text{travel}(a(y), y) \rangle$

STN: Methods (Partially-ordered)

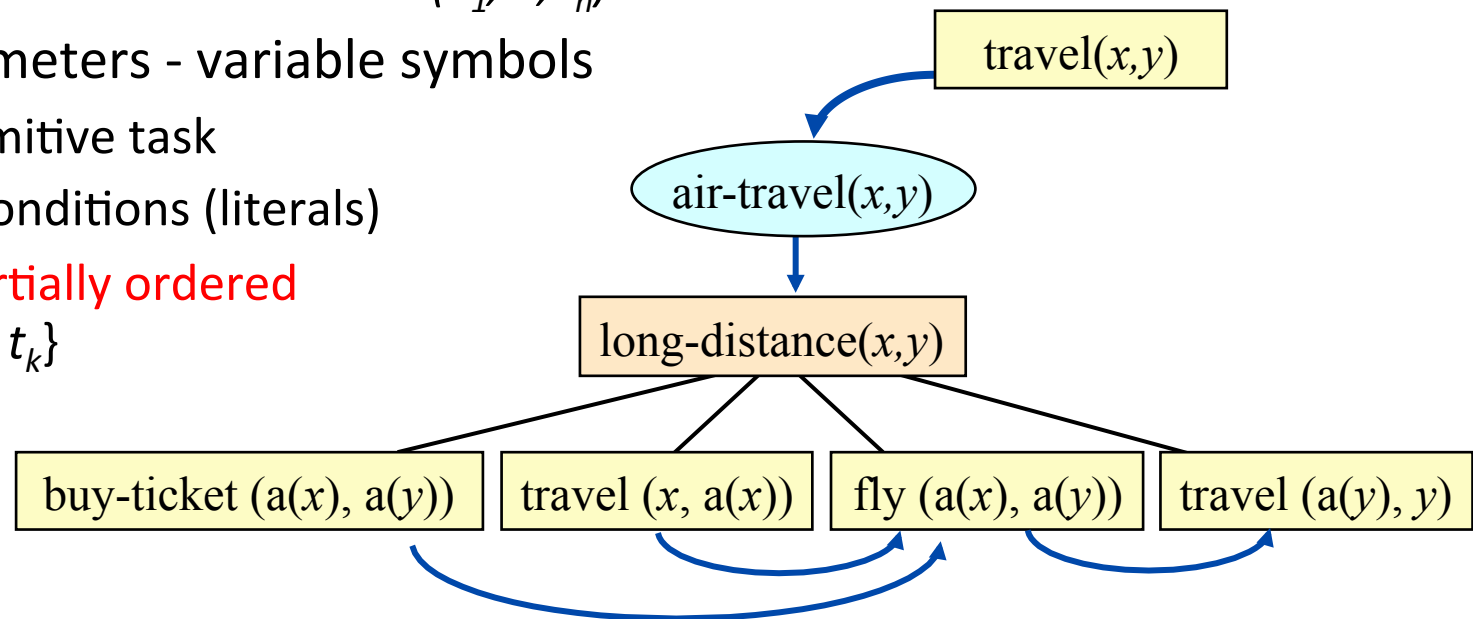
- Partially-ordered method: a 4-tuple
 $m = (\text{name}(m), \text{task}(m), \text{precond}(m), \text{subtasks}(m))$
 - $\text{name}(m)$: an expression of the form $n(x_1, \dots, x_n)$
 - x_1, \dots, x_n are parameters - variable symbols
 - $\text{task}(m)$: a nonprimitive task
 - $\text{precond}(m)$: preconditions (literals)
 - $\text{subtasks}(m)$: a **partially ordered** set of tasks $\{t_1, \dots, t_k\}$

air-travel(x,y)

task: travel(x,y)



precond: long-distance(x,y)

network: $u_1 = \text{buy-ticket}(a(x), a(y))$, $u_2 = \text{travel}(x, a(x))$, $u_3 = \text{fly}(a(x), a(y))$,
 $u_4 = \text{travel}(a(y), y)$, $\{(u_1, u_3), (u_2, u_3), (u_3, u_4)\}$




Example: DWR

- Task: Move three stacks of containers in a way that preserves the order of the containers
- One way to move each stack:
- First move the containers from p to an intermediate pile r
- Then move them from r to q



Example: DWR

- (informal) methods:
 - move each stack twice: move stack to intermediate pile (reversing order) and then to final destination (reversing order again)
 - move stack: repeatedly/recursively move the topmost container until the stack is empty
 - move top-most: take followed by put action



Example: DWR Total-Order Formulation

take-and-put($c, k, l_1, l_2, p_1, p_2, x_1, x_2$):

task: move-topmost-container(p_1, p_2)

precond: $\text{top}(c, p_1), \text{on}(c, x_1),$; true if p_1 is not empty
 $\text{attached}(p_1, l_1), \text{belong}(k, l_1),$; bind l_1 and k
 $\text{attached}(p_2, l_2), \text{top}(x_2, p_2)$; bind l_2 and x_2

subtasks: $\langle \text{take}(k, l_1, c, x_1, p_1), \text{put}(k, l_2, c, x_2, p_2) \rangle$

recursive-move(p, q, c, x):

task: move-stack(p, q)

precond: $\text{top}(c, p), \text{on}(c, x)$; true if p is not empty

subtasks: $\langle \text{move-topmost-container}(p, q), \text{move-stack}(p, q) \rangle$
 ;; the second subtask recursively moves the rest of the stack

do-nothing(p, q)

task: move-stack(p, q)

precond: $\text{top}(\text{pallet}, p)$; true if p is empty



subtasks: $\langle \rangle$; no subtasks, because we are done

move-each-twice()

task: move-all-stacks()

precond: ; no preconditions

subtasks: ; move each stack twice:
 $\langle \text{move-stack}(p1a, p1b), \text{move-stack}(p1b, p1c),$
 $\text{move-stack}(p2a, p2b), \text{move-stack}(p2b, p2c),$
 $\text{move-stack}(p3a, p3b), \text{move-stack}(p3b, p3c) \rangle$



STN: Solving Total-Order Planning Problems

TFD($s, \langle t_1, \dots, t_k \rangle, O, M$)

if $k = 0$ then return $\langle \rangle$ (i.e., the empty plan)

if t_1 is primitive then

$active \leftarrow \{(a, \sigma) \mid a \text{ is a ground instance of an operator in } O,$
 $\sigma \text{ is a substitution such that } a \text{ is relevant for } \sigma(t_1),$
 $\text{and } a \text{ is applicable to } s\}$

if $active = \emptyset$ then return failure

nondeterministically choose any $(a, \sigma) \in active$

$\pi \leftarrow \text{TFD}(\gamma(s, a), \sigma(\langle t_2, \dots, t_k \rangle), O, M)$

if $\pi = \text{failure}$ then return failure

else return $a.\pi$

else if t_1 is nonprimitive then

$active \leftarrow \{m \mid m \text{ is a ground instance of a method in } M,$
 $\sigma \text{ is a substitution such that } m \text{ is relevant for } \sigma(t_1),$
 $\text{and } m \text{ is applicable to } s\}$

if $active = \emptyset$ then return failure

nondeterministically choose any $(m, \sigma) \in active$

$w \leftarrow \text{subtasks}(m).\sigma(\langle t_2, \dots, t_k \rangle)$

return $\text{TFD}(s, w, O, M)$

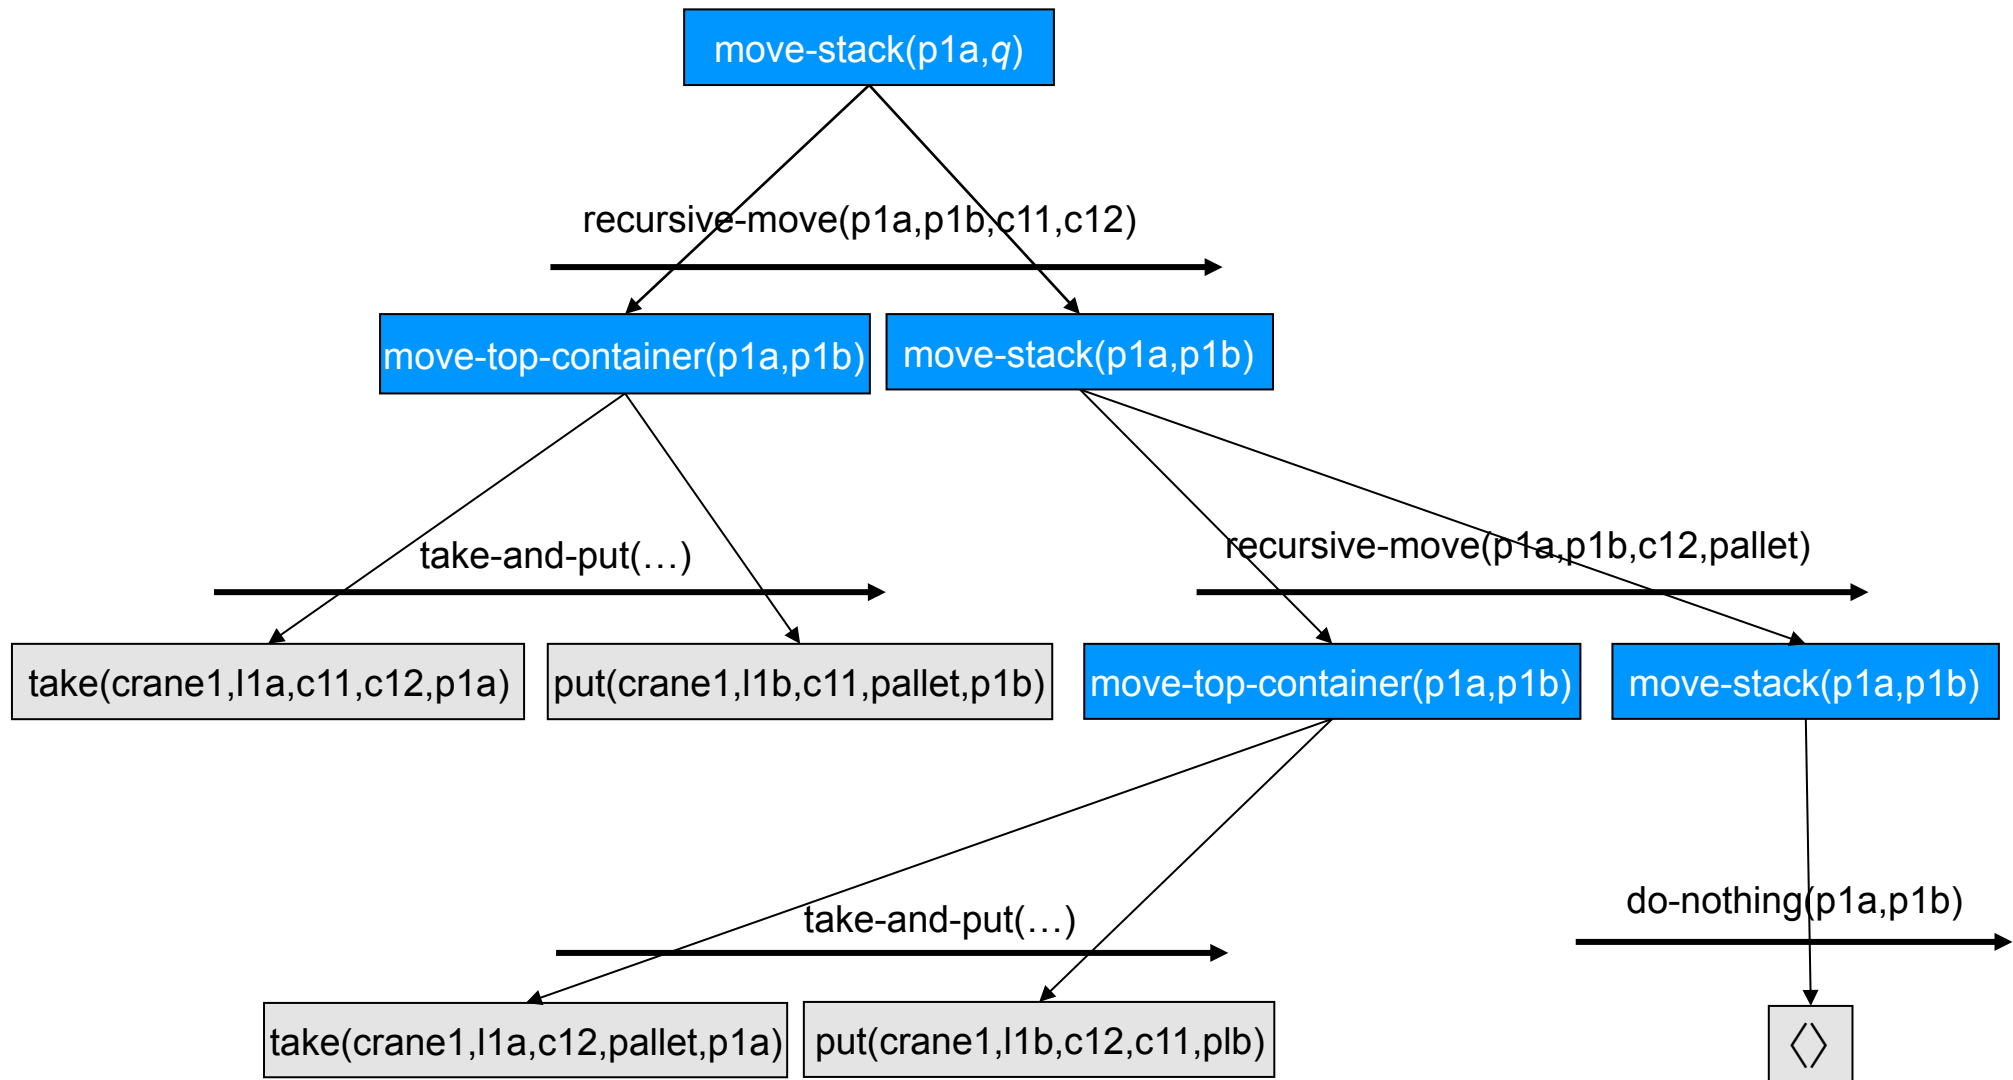
state s ; task list $T = (\mathbf{t_1}, t_2, \dots)$
 action a

state $\gamma(s, a)$; task list $T = (t_2, \dots)$

task list $T = (\mathbf{t_1}, t_2, \dots)$
 method instance m


task list $T = (\mathbf{u_1, \dots, u_k}, t_2, \dots)$

Example: DWR Decomposition Tree - TFD

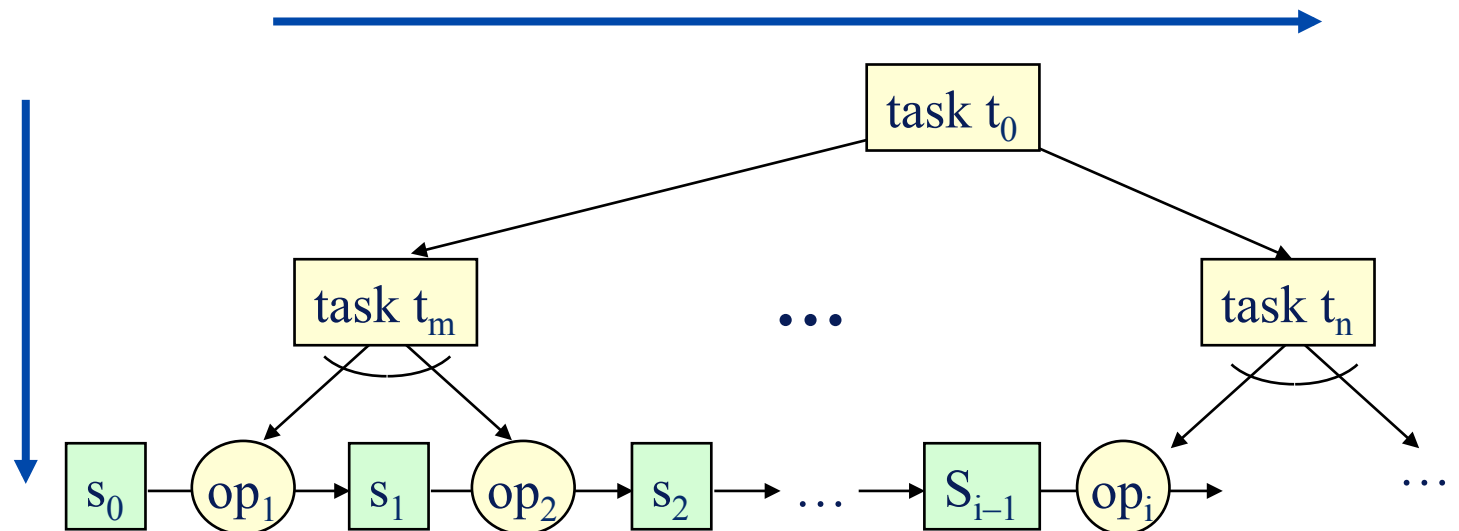


Comparison to Forward and Backward Search

- In state-space planning, must choose whether to search forward or backward



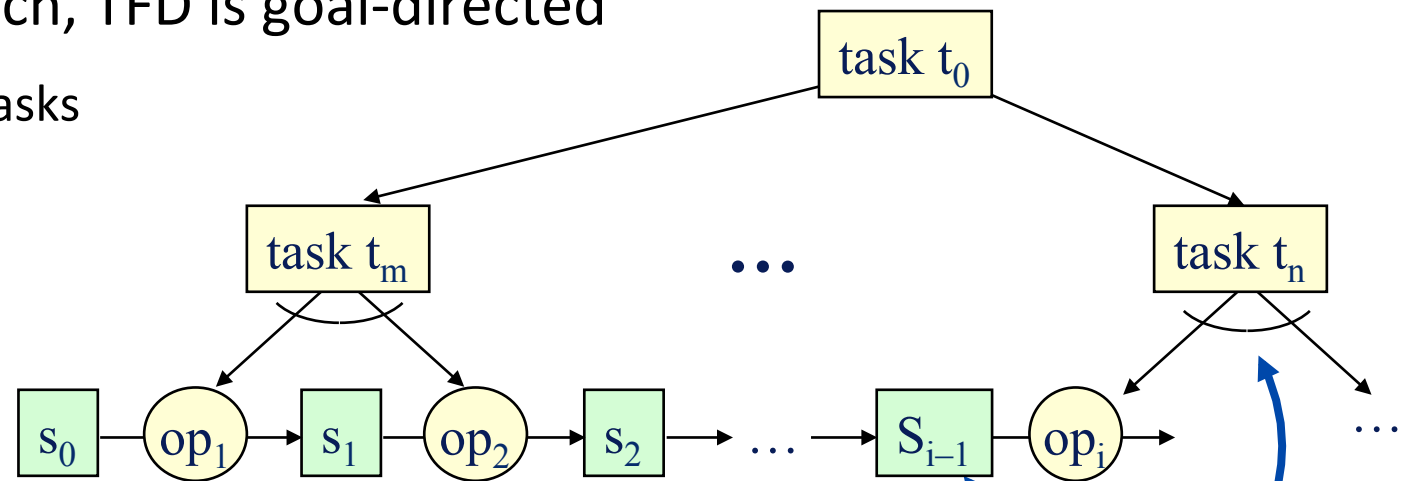
- In HTN planning, there are *two* choices to make about direction:
 - forward or backward
 - up or down



- TFD goes down and forward

Comparison to Forward and Backward Search

- Like a backward search, TFD is goal-directed
 - Goals correspond to tasks




- Like a forward search, it generates actions in the same order in which they'll be executed
- Whenever we want to plan the next task
 - we've already planned everything that comes before it
 - Thus, we know the current state of the world

Expressivity Relative to Classical Planning

- Any classical planning problem can be translated into an ordered-task planning problem in polynomial time
- Several ways to do this. One is roughly as follows:
 - For each goal or precondition e , create a task t_e
 - For each operator o and effect e , create a method $m_{o,e}$
 - Task: t_e
 - Subtasks: $t_{c1}, t_{c2}, \dots, t_{cn}, o$, where $c1, c2, \dots, cn$ are the preconditions of o
 - Partial-ordering constraints: each t_{ci} precedes o
- There are HTN planning problems that cannot be translated into classical planning problems at all
- Example on the next page

Example: Classical Planning can not represent this

- Two methods:
 - No arguments
 - No preconditions
- Two operators, a and b
 - Again, no arguments and no preconditions
- Initial state is empty, initial task is t
- Set of solutions is $\{a^n b^n \mid n > 0\}$
- No classical planning problem has this set of solutions
 - The state transition system is a finite state automaton
 - No finite state automaton can recognize $\{a^n b^n \mid n > 0\}$




Increasing Expressivity Further

- Knowing the current state makes it easy to do things that would be difficult otherwise
 - States can be arbitrary data structures

Us: East declarer, West dummy
Opponents: defenders, South & North
Contract: East – 3NT
On lead: West at trick 3

East: ♠KJ74
West: ♠A2
Out: ♠QT98653



- Preconditions and effects can include
 - logical inferences (e.g., Horn clauses)
 - complex numeric computations
 - interactions with other software packages
- Example: SHOP <http://www.cs.umd.edu/projects/shop>

SHOP (Simple Hierarchical Ordered Planner)

- Domain-independent algorithm for ordered task decomposition
 - Sound and complete
- Input:
 - State: a set of ground atoms
 - Task List: a linear list of tasks
 - Domain: methods, operators, axioms
- Output: one or more plans, it can return:
 - the first plan it finds
 - all possible plans
 - a least-cost plan
 - all least-cost plans

Example: SHOP

- **Initial task list:** ((travel home park))
- **Initial state:** ((at home) (cash 20) (distance home park 8))
- **Methods** (task, preconditions, subtasks):
 - (:method (travel ?x ?y)
((at x) (walking-distance ?x ?y)) ' (!walk ?x ?y)) 1)
 - (:method (travel ?x ?y)
((at ?x) (have-taxi-fare ?x ?y))
' (!call-taxi ?x) (!ride ?x ?y) (!pay-driver ?x ?y)) 1)
- **Axioms:**
 - (:- (walking-dist ?x ?y) ((distance ?x ?y ?d) (eval (<= ?d 5))))
 - (:- (have-taxi-fare ?x ?y)
((have-cash ?c) (distance ?x ?y ?d) (eval (>= ?c (+ 1.50 ?d)))))
- **Primitive operators** (task, delete list, add list)
 - (:operator (!walk ?x ?y) ((at ?x)) ((at ?y)))
 - ...

Optional cost;
default is 1

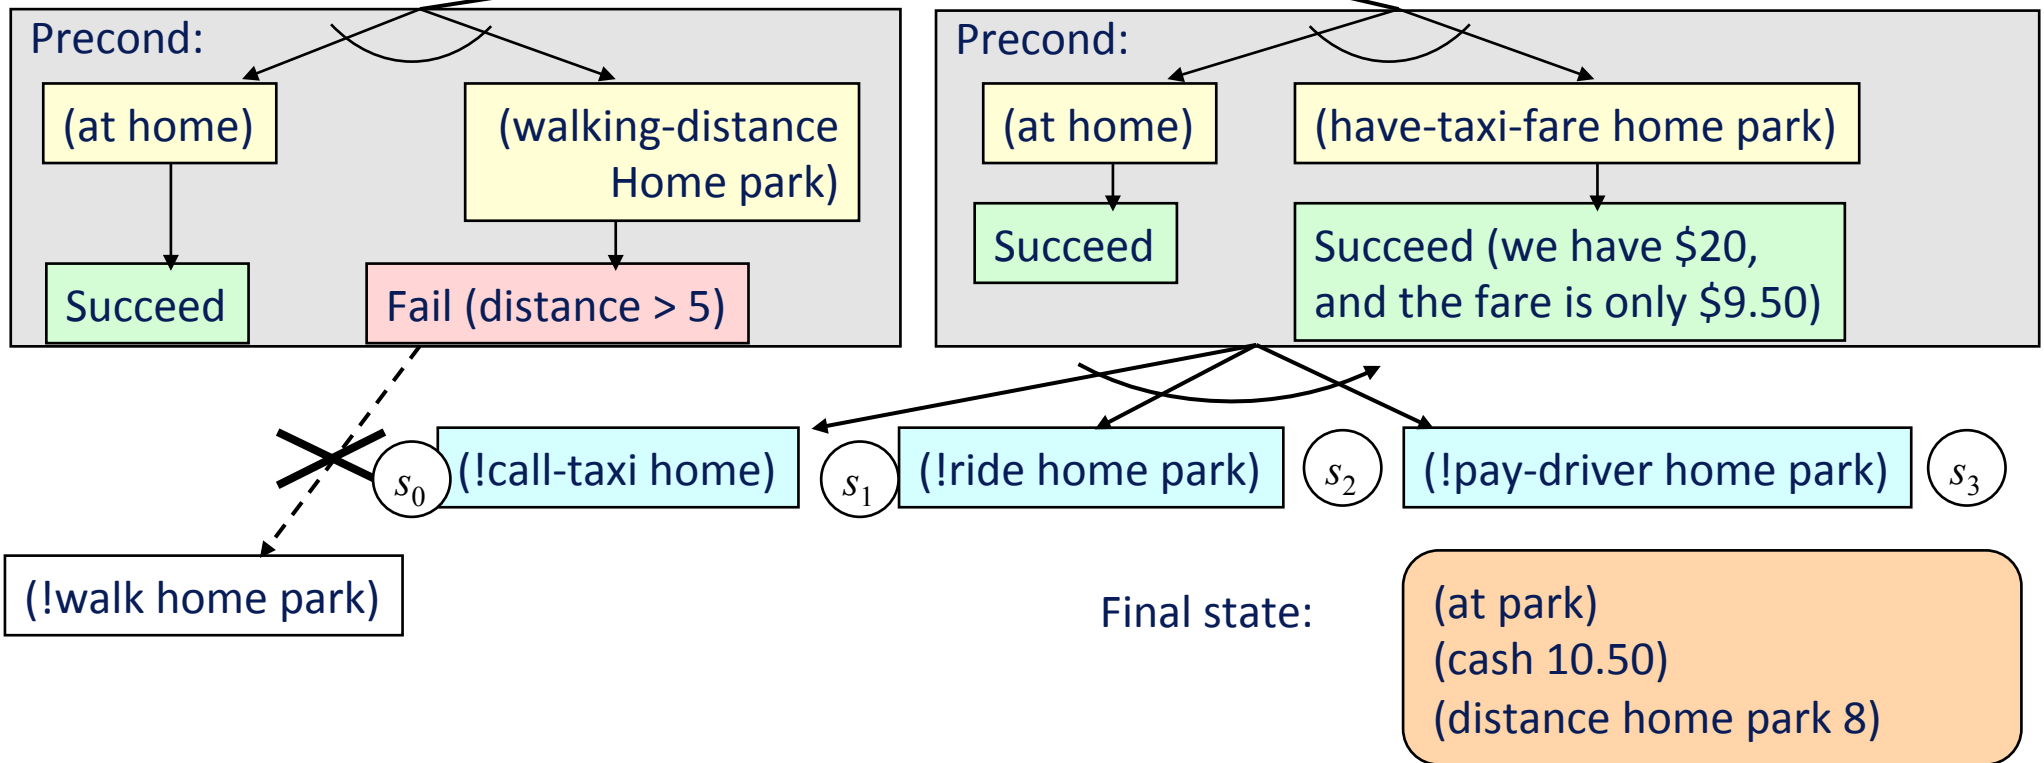


Example: SHOP (Continued)

Initial state:

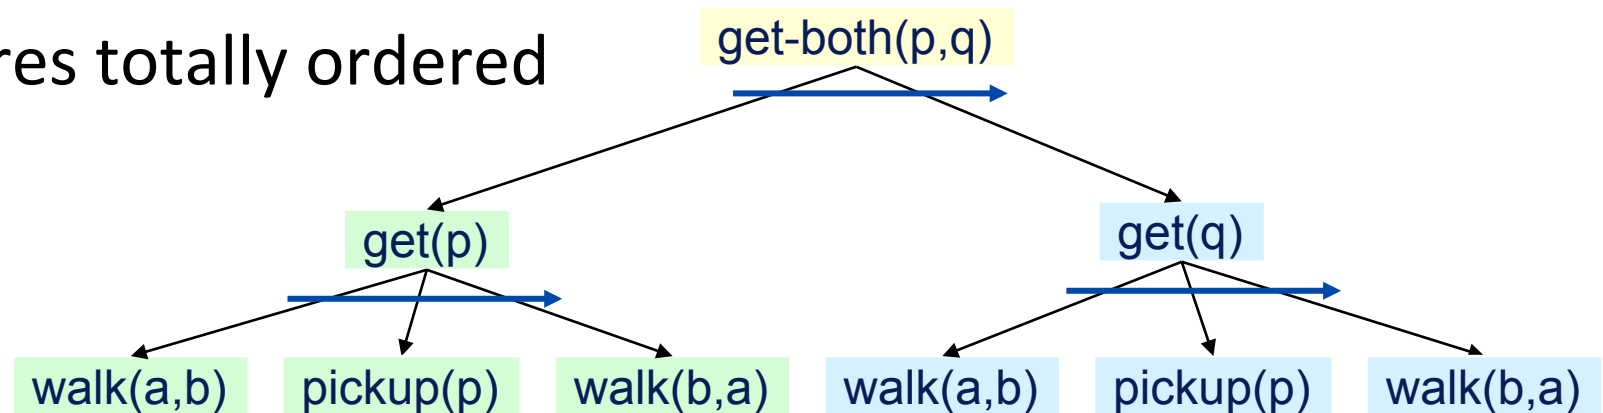
(at home)
(cash 20)
(distance home park 8)

(travel home park)

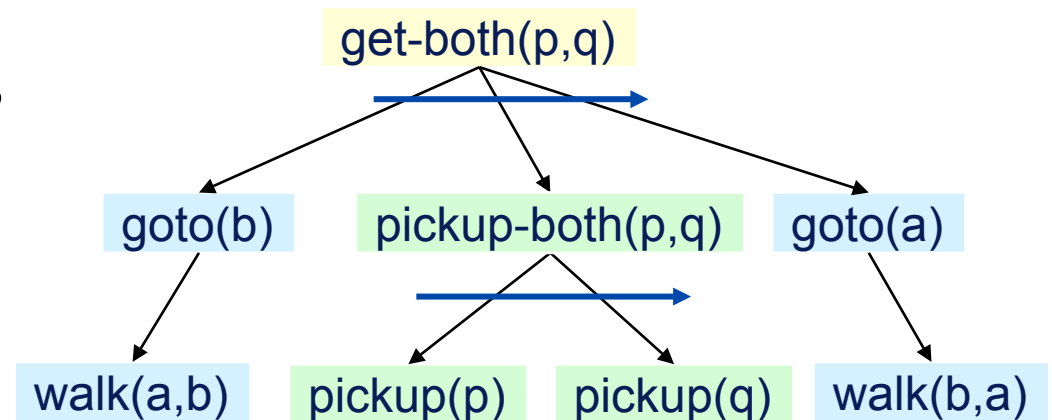


Limitation of Ordered-Task Planning

- TFD requires totally ordered methods

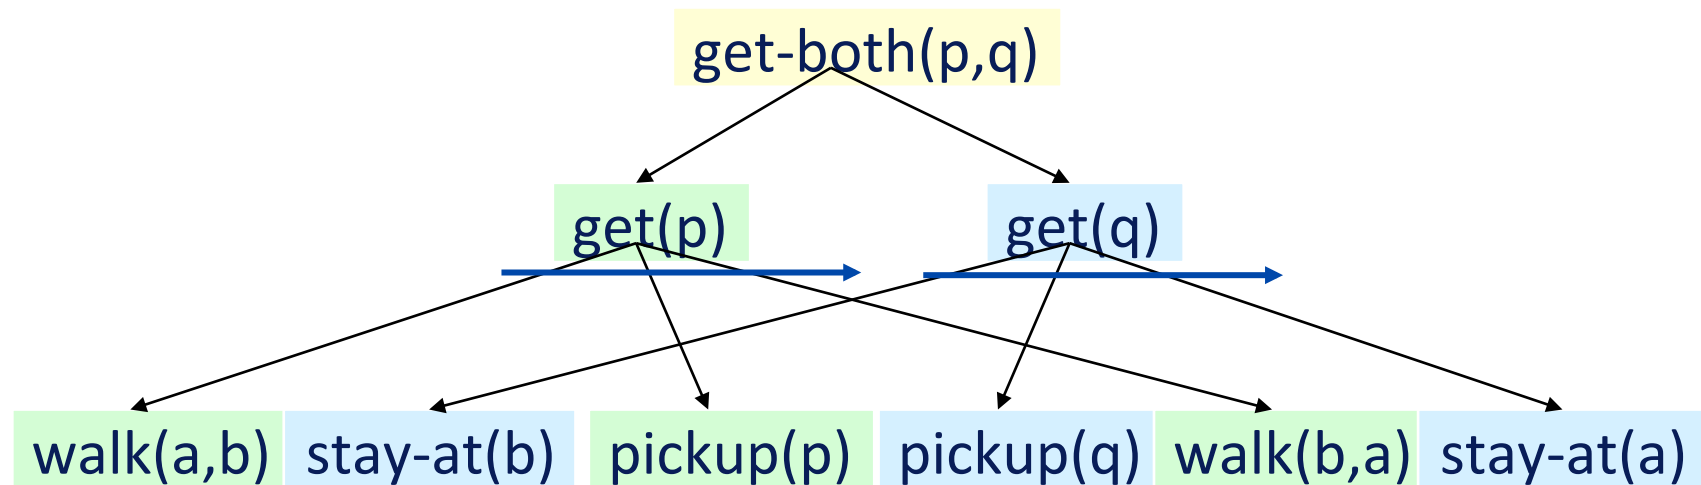


- Can't interleave subtasks of different tasks
- Sometimes this makes things awkward
 - Need to write methods that reason globally instead of locally



Generalize TFD to interleave subtasks

- Generalize methods to allow the subtasks to be partially ordered
- Consequence: plans may interleave subtasks of different tasks



- This makes the planning algorithm more complicated

Example: DWR Partial-Order Formulation

take-and-put($c, k, l_1, l_2, p_1, p_2, x_1, x_2$):

task: move-topmost-container(p_1, p_2)

precond: $\text{top}(c, p_1), \text{on}(c, x_1),$; true if p_1 is not empty
 $\text{attached}(p_1, l_1), \text{belong}(k, l_1),$; bind l_1 and k
 $\text{attached}(p_2, l_2), \text{top}(x_2, p_2)$; bind l_2 and x_2

subtasks: $\langle \text{take}(k, l_1, c, x_1, p_1), \text{put}(k, l_2, c, x_2, p_2) \rangle$

recursive-move(p, q, c, x):

task: move-stack(p, q)

precond: $\text{top}(c, p), \text{on}(c, x)$; true if p is not empty

subtasks: $\langle \text{move-topmost-container}(p, q), \text{move-stack}(p, q) \rangle$
 ;; the second subtask recursively moves the rest of the stack

do-nothing(p, q)

task: move-stack(p, q)

precond: $\text{top}(\text{pallet}, p)$; true if p is empty

subtasks: $\langle \rangle$; no subtasks, because we are done

move-each-twice()

task: move-all-stacks()

precond: ; no preconditions



network: ; move each stack twice:

$u_1 = \text{move-stack}(p1a, p1b), u_2 = \text{move-stack}(p1b, p1c),$

$u_3 = \text{move-stack}(p2a, p2b), u_4 = \text{move-stack}(p2b, p2c),$

$u_5 = \text{move-stack}(p3a, p3b), u_6 = \text{move-stack}(p3b, p3c),$

$\{(u_1, u_2), (u_3, u_4), (u_5, u_6)\}$



Solving Partial-Order STNs

$\text{PFD}(s, w, O, M)$

if $w = \emptyset$ then return the empty plan

nondeterministically choose any $u \in w$ that has no predecessors in w

if t_u is a primitive task then

$active \leftarrow \{(a, \sigma) \mid a \text{ is a ground instance of an operator in } O,$
 $\sigma \text{ is a substitution such that } \text{name}(a) = \sigma(t_u),$
 $\text{and } a \text{ is applicable to } s\}$

if $active = \emptyset$ then return failure

nondeterministically choose any $(a, \sigma) \in active$

$\pi \leftarrow \text{PFD}(\gamma(s, a), \sigma(w - \{u\}), O, M)$

if $\pi = \text{failure}$ then return failure

else return $a.\pi$

else

$active \leftarrow \{(m, \sigma) \mid m \text{ is a ground instance of a method in } M,$
 $\sigma \text{ is a substitution such that } \text{name}(m) = \sigma(t_u),$
 $\text{and } m \text{ is applicable to } s\}$

if $active = \emptyset$ then return failure

nondeterministically choose any $(m, \sigma) \in active$

nondeterministically choose any task network $w' \in \delta(w, u, m, \sigma)$

return($\text{PFD}(s, w', O, M)$)

$\pi = \{a_1, \dots, a_k\}; w = \{t_1, t_2, t_3, \dots\}$
 operator instance **a**

$\pi = \{a_1, \dots, a_k, \mathbf{a}\}; w' = \{t_2, t_3, \dots\}$

$w = \{t_1, t_2, \dots\}$
 method instance **m**

$w' = \{\mathbf{u}_1, \dots, \mathbf{u}_k, t_2, \dots\}$

Solving Partial-order STNs

PFD(s, w, O, M)

if $w = \emptyset$ then return the empty plan

nondeterministically choose any $u \in w$ that has no predecessors in w

- Intuitively, w is a partially ordered set of tasks $\{t_1, t_2, \dots\}$
 - But w may contain a task more than once
 - e.g., travel from UMD to LAAS twice
 - The mathematical definition of a set doesn't allow this
- Define w as a partially ordered set of *task nodes* $\{u_1, u_2, \dots\}$
 - Each task node u corresponds to a task t_u
- In my explanations, I talk about t and ignore u

else

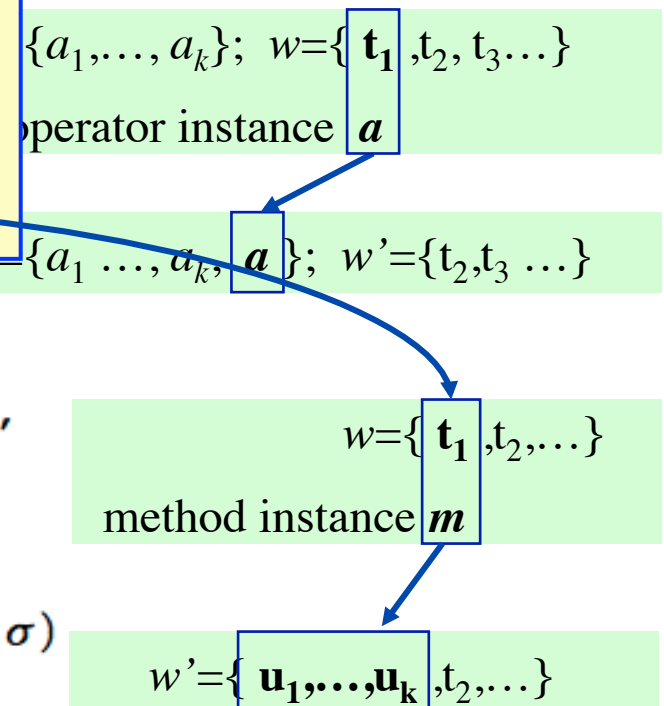
$active \leftarrow \{(m, \sigma) \mid m \text{ is a ground instance of a method in } M, \\ \sigma \text{ is a substitution such that } name(m) = \sigma(t_u), \\ \text{and } m \text{ is applicable to } s\}$

if $active = \emptyset$ then return failure

nondeterministically choose any $(m, \sigma) \in active$

nondeterministically choose any task network $w' \in \delta(w, u, m, \sigma)$

return(PFD(s, w', O, M))



Solving Partial-order STNs

PFD(s, w, O, M)

if $w = \emptyset$ then return the empty plan

nondeterministically choose any $u \in w$ that has no predecessors in w

if t_u is a primitive task then

$active \leftarrow \{(a, \sigma) \mid a \text{ is a ground instance of an operator in } O,$
 $\sigma \text{ is a substitution such that } name(a) = \sigma(t_u),$
 $\text{and } a \text{ is applicable to } s\}$

if $active = \emptyset$ then return failure

nondeterministically choose any $(a, \sigma) \in active$

$\pi \leftarrow \text{PFD}(\gamma(s, a), \sigma(w - \{u\}), O, M)$

if $\pi = \text{failure}$ then return failure

else return $a.\pi$

else

$active \leftarrow \{(m, \sigma) \mid m \text{ is a ground instance of a method in } M,$
 $\sigma \text{ is a substitution such that } name(m) = \sigma(t_u),$
 $\text{and } m \text{ is applicable to } s\}$

if $active = \emptyset$ then return failure

nondeterministically choose any $(m, \sigma) \in active$

nondeterministically choose any task network $w' \in \delta(w, u, m, \sigma)$

return(PFD(s, w', O, M))

$\pi = \{a_1, \dots, a_k\}; w = \{t_1, t_2, t_3, \dots\}$
 operator instance **a**

$\pi = \{a_1, \dots, a_k, \mathbf{a}\}; w' = \{t_2, t_3, \dots\}$

$w = \{t_1, t_2, \dots\}$
 method instance **m**

$w' = \{\mathbf{u}_1, \dots, \mathbf{u}_k, t_2, \dots\}$

Solving Partial-order STNs

PFD(s, w, O, M)

if $w = \emptyset$ then

nondeterministically choose any t_u

if t_u is a primitive task

$active \leftarrow \{t_u\}$

if $active = \emptyset$ then

nondeterministically choose any π

$\pi \leftarrow PFD(s, w, O, M)$

if $\pi = failure$ then

return π

else

$active \leftarrow \{(m, \sigma) \mid m \text{ is a ground instance of a method in } M, \sigma \text{ is a substitution such that } name(m) = \sigma(t_u), \text{ and } m \text{ is applicable to } s\}$

if $active = \emptyset$ then return failure

nondeterministically choose any $(m, \sigma) \in active$

nondeterministically choose any task network $w' \in \delta(w, u, m, \sigma)$

return(PFD(s, w', O, M))

$\delta(w, u, m, \sigma)$ has a complicated definition in the book. Here's what it means:

- We non-deterministically selected t_1 as the task to do first
- Must do t_1 's first subtask before the first subtask of every $t_i \neq t_1$
- Insert ordering constraints to ensure that this happens

t_2, t_3, \dots

$\pi = \{a_1, \dots, a_k, a\}; w' = \{t_2, t_3, \dots\}$

$w = \{t_1, t_2, \dots\}$

method instance m

$w' = \{u_1, \dots, u_k, t_2, \dots\}$

STN Summary

- PFD is sound and complete
- STN – simplified version of HTN
 - TFD – Total-order Forward Decomposition (used in SHOP)
 - Input: tasks are totally ordered
 - Output: totally ordered plan
 - PFD – Partial-order Forward Decomposition (SHOP2)
 - Input: tasks are partially ordered
 - Output: totally ordered plan
- SHOP2:
 - Won one of the top four awards in the AIPS-2002 Planning Competition
 - Freeware, open source
 - Implementation available at <http://www.cs.umd.edu/projects/shop>

STN v HTN

- HTN – generalization of STN
 - More freedom about how to construct the task networks.
 - Can use other decomposition procedures not just forward-decomposition.
 - Like Partial-Order Planning combined with STN
 - Input: Partial-order tasks
 - Output: The resulting plan is partially ordered
 - Plans can be totally ordered or partially ordered
 - Can have constraints associated with tasks and methods
 - Things that must be true before a state, in between two given states, or after a state (replaces STN preconditions)
 - Some algorithms use causal links and threats like those in PSP


TLPlan's Expressivity Compared with SHOP and SHOP2

- Equivalent expressive power
- Both know the current state at each step of the planning process, and use this to prune operators
- Both can call external subroutines
 - SHOP uses “eval” to call LISP functions
 - In TLPlan, a function symbol can correspond to a computed function
- Main difference
 - in SHOP and SHOP2, the methods talk about what *can* be done
 - SHOP and SHOP2 don't do anything unless a method says to do it
 - TLPlan's control rules talk about what *cannot* be done
 - TLPlan will try everything that the control rules don't prohibit
- Which approach is more convenient depends on the problem domain

Experimental Comparison


- Several years ago, we did a comparison of SHOP, TLPlan, and Blackbox
 - Blackbox is a domain-independent planner that uses a combination of Graphplan and satisfiability
 - One of the two fastest planners in the 1998 planning competition
- Test domain: the logistics domain
 - A classical planning problem
 - Much simpler than real logistics planning
 - Scenario: use trucks and airplanes to deliver packages
 - Like a simplified version of the DWR domain in which containers don't get stacked on each other
- Test conditions
 - SHOP and TLPlan on a 167-MHz Sun Ultra with 64 MB of RAM
 - We couldn't run Blackbox on our machine
 - Published results: Blackbox on a faster machine with 8 GB of RAM

Logistics Domain Results




Logistics Domain Results (continued)

- Same graph as before, but on a logarithmic scale



| Average | Blackbox | TLPlan | SHOP |
|----------|----------|--------|------|
| CPU time | 327.1 | 2.9 | 1.1 |

- Number of actions in the plans



| Average no. | Blackbox | TLPlan | SHOP |
|-------------|----------|--------|------|
| of actions | 82.5 | 54.5 | 51.9 |

Summary: Results

- TLPlan and SHOP took similar amounts of time
 - In this experiment, SHOP was slightly faster, but in others TLPlan may be faster
- Blackbox took about 1000 times as much time and needed about 100 times as much memory
- Reasons why:
 - SHOP's input included domain-specific methods & axioms
 - TLPlan's input included domain-specific control rules
 - This enabled them to find near-optimal solutions in low-order polynomial time and space
 - Blackbox is a fully automated planner
 - No domain-specific knowledge
 - trial-and-error search, exponential time and space

Domain-Configurable Planners

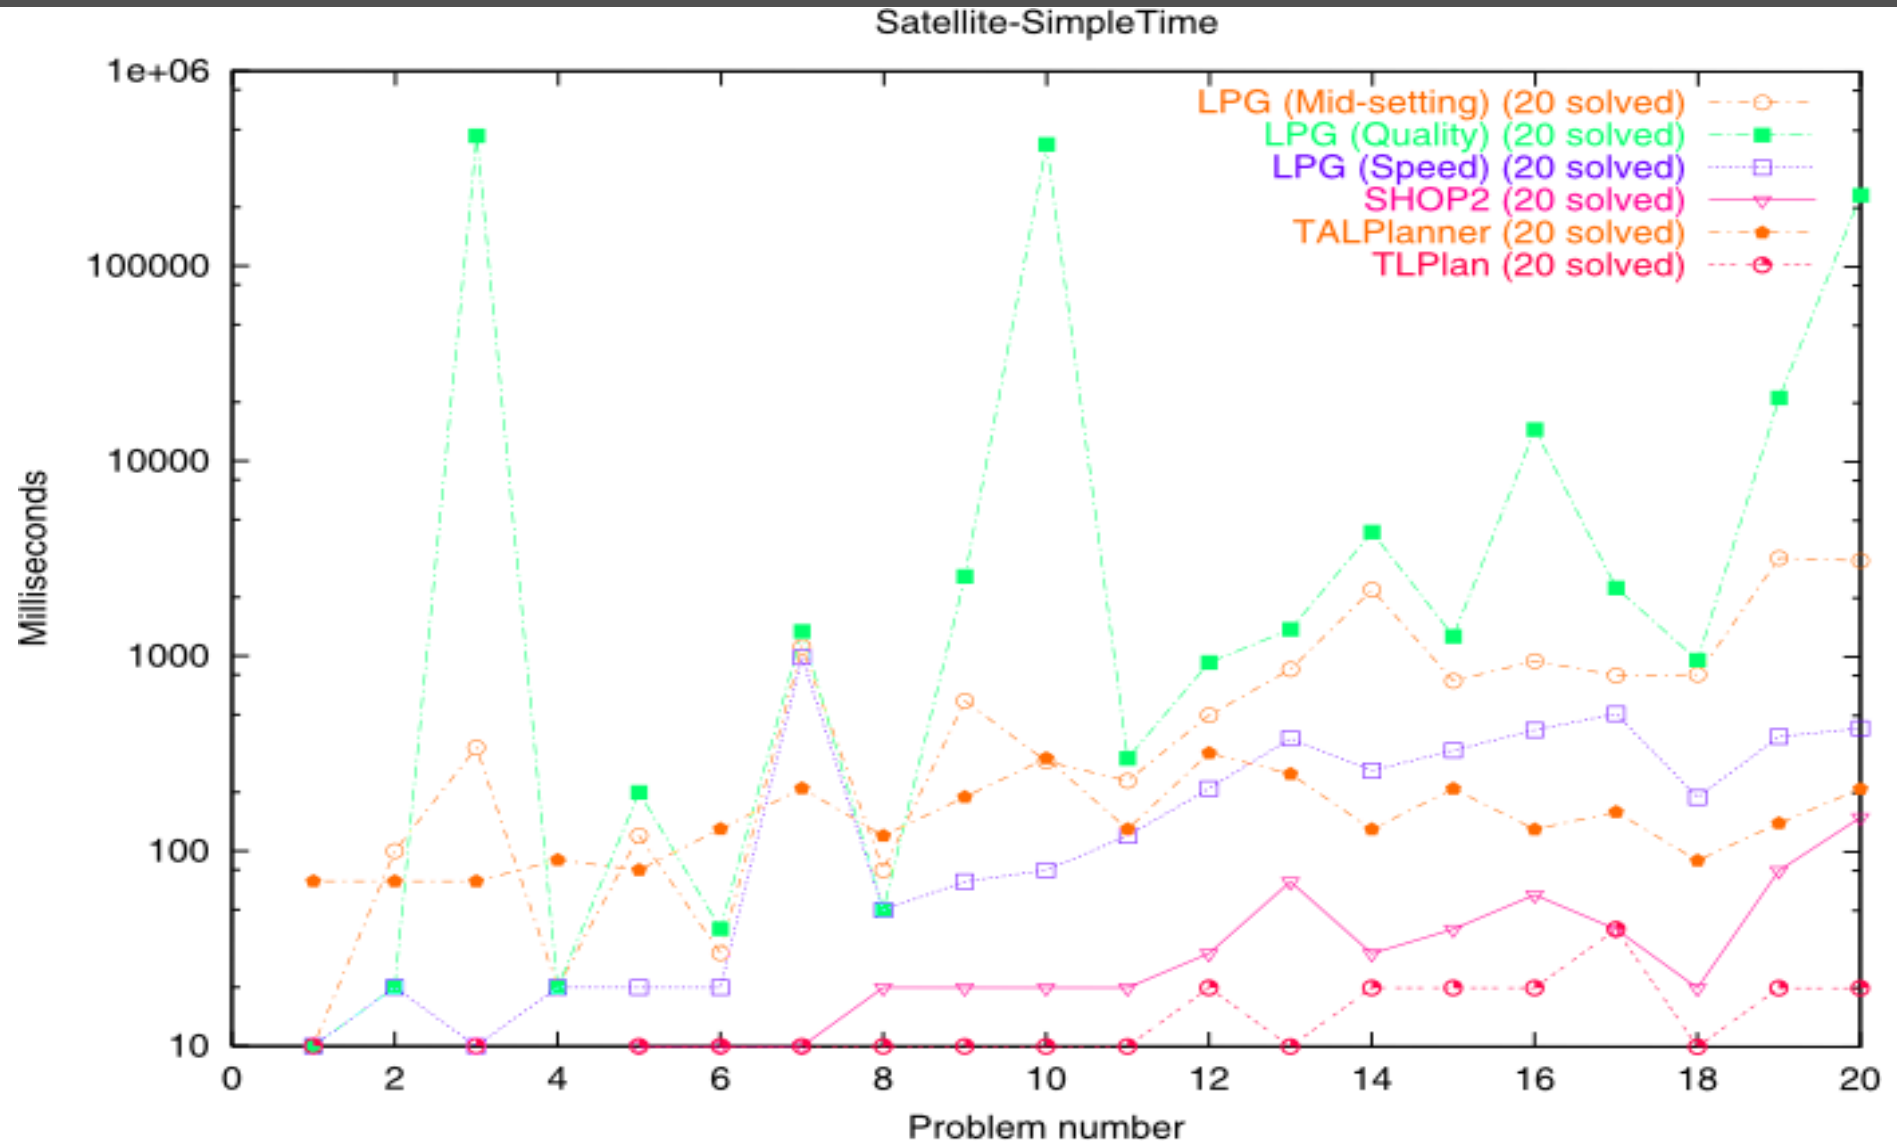
Compared to Classical Planners

- Disadvantage: writing a knowledge base can be more complicated than just writing classical operators
- Advantage: can encode “recipes” as collections of methods and operators
 - Express things that can’t be expressed in classical planning
 - Specify standard ways of solving problems
 - Otherwise, the planning system would have to derive these again and again from “first principles,” every time it solves a problem
 - Can speed up planning by many orders of magnitude (e.g., polynomial time versus exponential time)


Example from the AIPS-2002 Competition

- The satellite domain
 - Planning and scheduling observation tasks among multiple satellites
 - Each satellite equipped in slightly different ways
- Several different versions. Results are shown for the following:
 - *Simple time*:
 - concurrent use of different satellites
 - data can be acquired more quickly if they are used efficiently
 - *Numeric*:
 - fuel costs for satellites to slew between targets; finite amount of fuel available.
 - data takes up space in a finite capacity data store
 - Plans are expected to acquire all the necessary data at minimum fuel cost.
 - *Hard Numeric*:
 - *no logical goals at all* – thus even the null plan is a solution
 - Plans that acquire more data are better – thus the null plan has no value
 - None of the classical planners could handle this


Satellite Problem Domain: Simple Time: Runtime



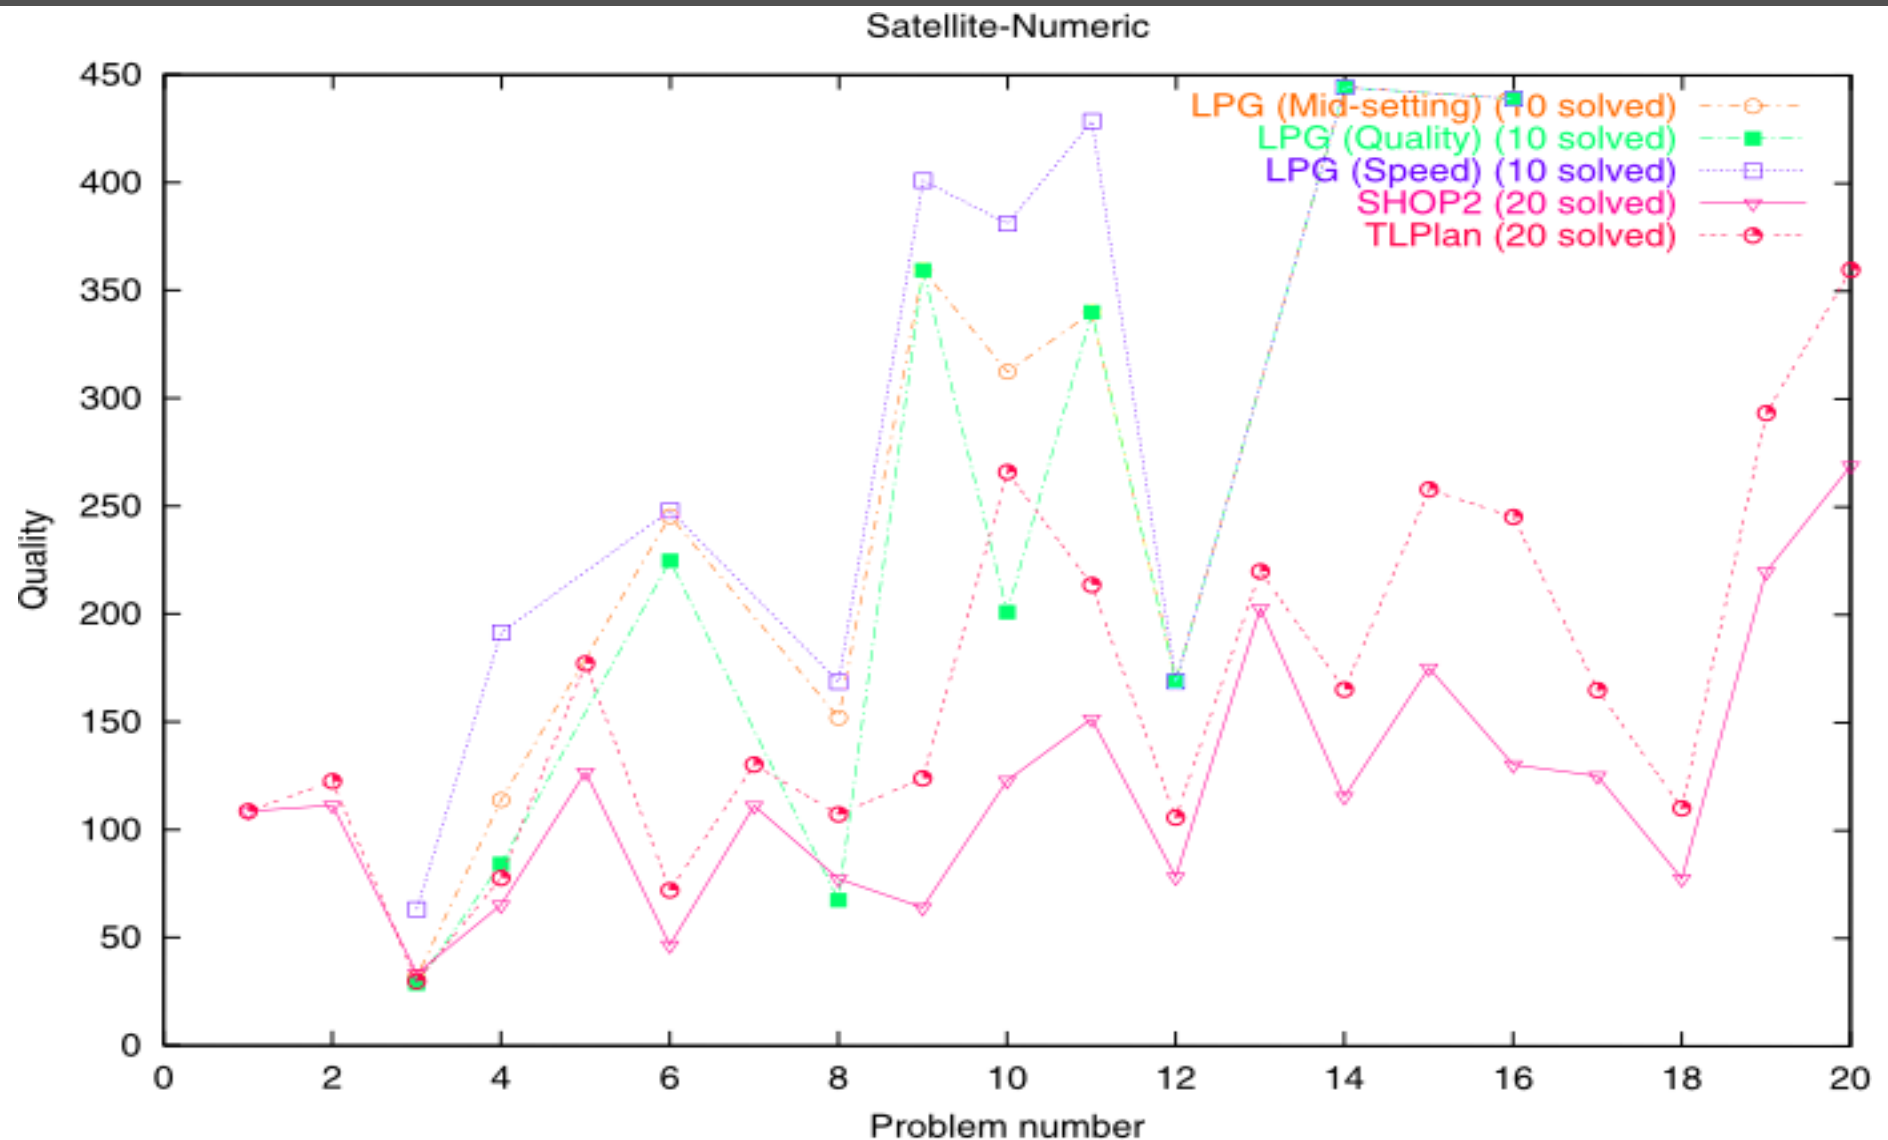
Satellite Problem Domain: Simple Time: Plan Quality




Satellite Problem Domain: Numeric: Runtime



Satellite Problem Domain: Numeric: Plan Quality



Satellite Problem Domain: Hard Numeric: Runtime



Satellite Problem Domain: Hard Numeric: Plan Quality

